
CS3402: Chapter 3

SQL: Structured Query Language

SQL: Structured Query Language

- SQL is a hybrid language, its basically 4 types of languages in one
 - ◆ Data Definition Language(DDL)
used for define database schemas
 - ◆ Data Query Language (DQL)
used to query the database for information
 - ◆ Data Manipulation Language (DML)
used for inserting, updating and deleting data from the database
 - ◆ Data Control Language (DCL)
used for controlling the access to the data in the database
- Each statement in SQL ends with a semicolon (;)
- SQL is case insensitive
- Implementations of SQL in Different DBMS (MySQL, Oracle) are slightly different

Data Definition Language(DDL)

CREATE Table

- Create Database:

 - ◆ CREATE SCHEMA COMPANY;

- Create TABLE:

 - ◆ CREATE TABLE EMPLOYEE ...

- Base tables (base relations)

 - ◆ Relation and its tuples are **actually** (physically) created and stored as a **file** by the DBMS

 - ◆ The tables are stored in the secondary storage in the **specified format**

- Virtual relations

 - ◆ Created through the CREATE VIEW statement

 - ◆ The tables are **not actually created** but are presented to the user through reconstruction (view) of base tables

 - ◆ View is always up-to-date (change in base table will be reflected in the views automatically)

Table Manipulation

- Table creation: `CREATE TABLE table_name`
 `(column_name1 data_type(size),`
 `column_name2 data_type(size),....);`
`CREATE VIEW view_name AS SELECT FROM...`
- Table deletion: `DROP TABLE table-name;`
 `DROP VIEW view-name`
- Table update:
 `ALTER TABLE table-name`
 `ADD Aj, Dj`
 (to add new attribute Aj with domain Dj to an existing table)
 `ALTER VUEW view-name`

One possible database state for the COMPANY relational database schema

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

One possible database state for the COMPANY relational database schema

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

PROJECT

<u>Pname</u>	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse

SQL CREATE TABLE data definition statements for defining the COMPANY schema

CREATE TABLE EMPLOYEE

(Fname VARCHAR(15)
Minit CHAR,
Lname VARCHAR(15)
Ssn CHAR(9)
Bdate DATE,
Address VARCHAR(30),
Sex CHAR,
Salary DECIMAL(10,2),
Super_ssn CHAR(9),
Dno INT

NOT NULL,

NOT NULL,
NOT NULL,

NOT NULL,

Indicate this column does not accept NULL value

Define primary key

• PRIMARY KEY (Ssn),

CREATE TABLE DEPARTMENT

(Dname VARCHAR(15)
Dnumber INT
Mgr_ssn CHAR(9)
Mgr_start_date DATE,

NOT NULL,
NOT NULL,
NOT NULL,

ensures that all values in a column are different

• PRIMARY KEY (Dnumber),

• UNIQUE (Dname),

FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn));

Define foreign key and its reference

CREATE TABLE DEPT_LOCATIONS

(Dnumber INT
Dlocation VARCHAR(15)

NOT NULL,
NOT NULL,

PRIMARY KEY (Dnumber, Dlocation),

FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber));

SQL CREATE TABLE data definition statements for defining the COMPANY schema

```
CREATE TABLE PROJECT
( Pname                VARCHAR(15)          NOT NULL,
  Pnumber              INT                  NOT NULL,
  Plocation            VARCHAR(15),
  Dnum                 INT                  NOT NULL,
  PRIMARY KEY (Pnumber),
  UNIQUE (Pname),
  FOREIGN KEY (Dnum) REFERENCES DEPARTMENT(Dnumber) );

CREATE TABLE WORKS_ON
( Essn                 CHAR(9)              NOT NULL,
  Pno                  INT                  NOT NULL,
  Hours                DECIMAL(3,1)         NOT NULL,
  PRIMARY KEY (Essn, Pno),
  FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn),
  FOREIGN KEY (Pno) REFERENCES PROJECT(Pnumber) );

CREATE TABLE DEPENDENT
( Essn                 CHAR(9)              NOT NULL,
  Dependent_name        VARCHAR(15)         NOT NULL,
  Sex                   CHAR,
  Bdate                 DATE,
  Relationship           VARCHAR(8),
  PRIMARY KEY (Essn, Dependent_name),
  FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn) );
```

Data Query Language(DQL)

The *SELECT-FROM-WHERE* Structure of Basic SQL Queries

■ *SELECT* statement

◆ The basic statement for retrieving information from a database

SELECT <attribute list>
FROM <table list>
WHERE <condition>;

where

SELECT <attribute list>
FROM <table list>
[**WHERE** <condition>]
[**ORDER BY** <attribute list>];

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

Specified attributes



Satisfy the
conditions



Title	Year	Length	Type
Star War	1977	124	Color
Mighty Duck	1991	104	Color
Wayne's World	1992	95	Color

Retrieval from a single table

Query 0. Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

Q0: **SELECT** Bdate, Address ← Projection attributes
 FROM EMPLOYEE
 WHERE Fname='John' AND Minit='B' AND Lname='Smith'; ← Selection conditions

<u>Bdate</u>	<u>Address</u>
1965-01-09	731 Fondren, Houston, TX

- ◆ Logic operators to connect multiple conditions: **Not, AND, OR**
- ◆ Priority: (), NOT, AND OR
- ◆ E.g. Lname = 'Lee' OR Lname = 'Smith' AND Sex = 'Male'
(Lname = 'Lee' OR Lname = 'Smith') AND Sex = 'Male'

Retrieval from two tables

Query 1. Retrieve the name and address of all employees who work for the 'Research' department.

Q1: **SELECT** Fname, Lname, Address
 FROM EMPLOYEE, DEPARTMENT ← Two tables
 WHERE Dname='Research' **AND** Dnumber=Dno; ← Join conditions

<u>Fname</u>	<u>Lname</u>	<u>Address</u>
John	Smith	731 Fondren, Houston, TX
Franklin	Wong	638 Voss, Houston, TX
Ramesh	Narayan	975 Fire Oak, Humble, TX
Joyce	English	5631 Rice, Houston, TX

Retrieval from two tables: Join Operation

A B					
ABC	DEF				

X
JOIN

C	D	E	F
XYZ			
AAAA			
BBBB			
CCCC			

A	B	C	D	E	F
ABC	DEF	XYZ			
ABC	DEF	AAAA			
ABC	DEF	BBBB			
ABC	DEF	CCCC			

Retrieval from two tables

(c)

<u>Pnumber</u>	<u>Dnum</u>	<u>Lname</u>	<u>Address</u>	<u>Bdate</u>
10	4	Wallace	291Berry, Bellaire, TX	1941-06-20
30	4	Wallace	291Berry, Bellaire, TX	1941-06-20

Query 2. For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.

Q2: **SELECT** Pnumber, Dnum, Lname, Address, Bdate
 FROM PROJECT, DEPARTMENT, EMPLOYEE
 WHERE Dnum=Dnumber AND Mgr_ssn=Ssn AND
 Plocation='Stafford';

← Select-project-join
query


Project joins
Department

Department joins
Employee

Ambiguous Attribute Names

- Same name can be used for two (or more) attributes in different relations
 - ◆ As long as the attributes are in different relations
 - ◆ Must **qualify** the attribute name with the relation name to prevent ambiguity

Q1A: **SELECT** Fname, EMPLOYEE.Name, Address
 FROM EMPLOYEE, DEPARTMENT
 WHERE DEPARTMENT.Name='Research' **AND**
 DEPARTMENT.Dnumber=EMPLOYEE.Dnumber;



Aliasing, and Renaming

- Aliases or tuple variables

- ◆ Declare alternative relation names E and S to refer to the EMPLOYEE relation twice in a query:

Query 8. For each employee, retrieve the employee's first and last name and the first and last name of his or her immediate supervisor.

```
SELECT E.Fname, E.Lname, S.Fname, S.Lname  
FROM EMPLOYEE AS E, EMPLOYEE AS S  
WHERE E.Super_ssn=S.Ssn;
```

One level
recursive query

Alias

- ◆ Recommended practice to abbreviate names and to prefix same or similar attribute from multiple tables

Aliasing, and Renaming

- ◆ The attribute names can also be renamed

```
EMPLOYEE AS E(Fn, Mi, Ln, Ssn, Bd, Addr, Sex,  
              Sal, Sssn, Dno)
```


- ◆ The “AS” may be dropped in most SQL implementations

```
EMPLOYEE E(Fn, Mi, Ln, Ssn, Bd, Addr, Sex,  
           Sal, Sssn, Dno)
```

DISTINCT, LIKE

- *Find the first name of the employees:*

```
SELECT DISTINCT Fname  
FROM EMPLOYEE;
```



Records with same
values will be
removed

- *Find the names of all employees whose first name has the substring 'mm' included*

```
SELECT Fname, Minit, Lname  
FROM EMPLOYEE  
WHERE Fname LIKE "%mm%";
```

(Note: if we use "__mm%", then it becomes a special case)



3rd character (case sensitive)

Ordering Tuples

- List first name of all employees in alphabetic order

```
SELECT DISTINCT Fname  
FROM EMPLOYEE  
ORDER BY Fname;
```

By default, in ascending order.

- List the employee names in descending order of last name, and if several employees have the same last name, order them in ascending order by the first name

```
SELECT Fname, Minit, Lname  
FROM EMPLOYEE  
ORDER BY Lname DESC, Fname ASC;
```

Unspecified WHERE Clause and Use of the Asterisk

- Missing WHERE clause
 - ◆ Indicates no condition on tuple selection (select ALL)
- Effect is a CROSS PRODUCT (JOIN $n \times m$)
 - ◆ Result is all possible tuple combinations result

Queries 9 and 10. Select all EMPLOYEE Ssns (Q9) and all combinations of EMPLOYEE Ssn and DEPARTMENT Dname (Q10) in the database.

Q9: SELECT Ssn
 FROM EMPLOYEE;

Q10: SELECT Ssn, Dname
 FROM EMPLOYEE, DEPARTMENT;

← May return many tuples
Join operation

Unspecified WHERE Clause and Use of the Asterisk

- Specify an asterisk (*)
 - ◆ Retrieve all the attribute values of the selected tuples

Q1C: SELECT *
 FROM EMPLOYEE
 WHERE Dno=5;

Q1D: SELECT *
 FROM EMPLOYEE, DEPARTMENT
 WHERE Dname='Research' AND Dno=Dnumber;

Q10A: SELECT *
 FROM EMPLOYEE, DEPARTMENT;

←
EMPLOYEE: 100 tuples
DEPARTMENT: 10 tuples
Finally, 1000 tuples and all attributes

Tables as Sets in SQL

- SQL does not automatically eliminate duplicate tuples (the attributes of two tuples have same values) in query results (**NOT** a set)
- Use the keyword **DISTINCT** in the **SELECT** clause
 - ◆ Only distinct tuples should remain in the result

Query 11. Retrieve the salary of every employee (Q11) and all distinct salary values (Q11A).

Q11: **SELECT** **ALL** Salary
 FROM EMPLOYEE;

Q11A: **SELECT** **DISTINCT** Salary
 FROM EMPLOYEE;

Tables as Sets in SQL

■ Set operations

◆ UNION, INTERSECT, MINUS/EXCEPT (difference),

Query 4. Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

Q4A: (SELECT
FROM
WHERE

UNION
(SELECT
FROM
WHERE

DISTINCT Pnumber
PROJECT, DEPARTMENT, EMPLOYEE
Dnum=Dnumber AND Mgr_ssn=Ssn
AND Lname='Smith')

DISTINCT Pnumber
PROJECT, WORKS_ON, EMPLOYEE
Pnumber=Pno AND Essn=Ssn
AND Lname='Smith');

Same attribute

The projects in
the dept with
"Smith" as
manager

Tables as Sets in SQL, UNION

- Assuming the following relation schemes:

Customer (cname, street, city)

Branch (bname, assets, b-city)

Borrow (bname, loan#, cname, amount)

Deposit (bname, acct#, cname, balance)

- ◆ *Find all customers of the Kowloon branch*

```
(SELECT cname  
FROM Deposit  
WHERE bname = "Kowloon")
```

} A set of tuples

UNION

```
(SELECT cname  
FROM Borrow  
WHERE bname = "Kowloon");
```

} A set of tuples

Tables as Sets in SQL, MINUS/EXCEPT

- Assuming the following relation schemes:

Customer (cname, street, city)

Branch (bname, assets, b-city)

Borrow (bname, loan#, cname, amount)

Deposit (bname, acct#, cname, balance)

- ◆ *Find all customers of who has deposit account in Kowloon branch but no loan account in Kowloon branch*

(SELECT cname
FROM Deposit
WHERE bname = "Kowloon")

} A set of tuples

MINUS

(SELECT cname
FROM Borrow
WHERE bname = "Kowloon");

} A set of tuples

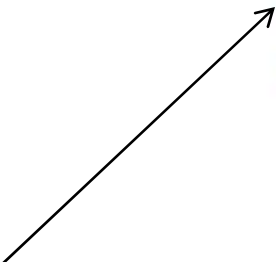
Nested Queries and Set / Multiset Comparisons

- Nested queries
 - ◆ select-from-where blocks within WHERE clause of another query
 - ◆ E.g., some queries require that existing values in the database be fetched and then used in a comparison condition
- Comparison operator `IN`
 - ◆ Compares value v with a set (or multiset) of values V
 - ◆ Evaluates to `TRUE` if v is one of the elements in V

Nested Queries, IN

- Use tuples of values in comparisons
 - ◆ Place them within parentheses

```
SELECT  DISTINCT Essn
FROM    WORKS_ON
WHERE   (Pno, Hours) IN ( SELECT  Pno, Hours
                          FROM    WORKS_ON
                          WHERE   Essn='123456789' );
```



Select the Essn of all employees who work the same (project, hours) as the employee Essn = "123456789"

Nested Query, IN

- Find all customers who have an account at some branch in which Jones has an account

solution 1:

```
SELECT DISTINCT T.cname
FROM Deposit T
WHERE T.cname != "Jones"
AND T.bname IN (SELECT S.bname
FROM Deposit S
WHERE S.cname = "Jones");
```

Deposit = T

Deposit = S

A set of records with "Jones" as cname →

solution 2:

```
SELECT DISTINCT T.cname
FROM Deposit S, Deposit T
WHERE S.cname = "Jones" AND S.bname = T.bname
AND T.cname != S.cname;
```

Nested Query, EXISTS

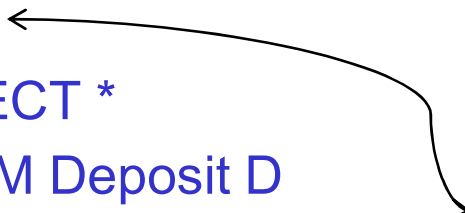
- EXISTS function
 - ◆ Check whether the result of a nested query is empty or not
 - ◆ Typically used in conjunction with a correlated nested query

- Correlated nested query
 - ◆ Evaluated once for each tuple in the outer query
 - ◆ Whenever a condition in the WHERE clause of a nested query references some attributes of a relation declared in the outer query

Nested Query, EXISTS

- Find all customers of Central branch who have an account there but no loan there

```
SELECT C.cname
FROM Customer C
WHERE EXISTS (SELECT *
              FROM Deposit D
              WHERE D.cname = C.cname
                 AND D.bname = "Central")
AND NOT EXISTS
  (SELECT *
   FROM Borrow B
   WHERE B.cname = C.cname
      AND B.bname = "Central");
```



Difference between EXISTS and IN

Query the record in t1 which has different value to any of the record in t2

solution 1: select * from table1 t1
 where not **exists** (select * from table2 t2
 where t2.value = t1.value);

solution 2: select *
 from table1 t1
 where t1.value not **in** (select t2.value
 from table2 t2);

Q1: Solution 1 and 2 are equivalent or not?

Q2: Which solution is more efficient?

Nested Queries, SOME/ANY, ALL

- Comparison operators that can be combined with ANY (or SOME)
ALL: = , >, >=, <, <=, and <>
- ◆ **ANY** (or **SOME**) operator
 - ◆ The **ANY** operator returns true if any of the subquery values meet the condition.
- ◆ **ALL** operator
 - ◆ The **ALL** operator returns true if all of the subquery values meet the condition.

Nested Queries, ALL

- Find names of all employees that have more salary than all employees in the department with department no. 5

```
SELECT  Lname, Fname
FROM    EMPLOYEE
WHERE   Salary > ALL ( SELECT  Salary
                        FROM    EMPLOYEE
                        WHERE   Dno=5 );
```

↖
The salary of Employee
tuples with Dno = 5


Nested Query, ALL

- Find branches having greater assets than all branches in N.T.

solution 1:

```
SELECT bname
FROM Branch
WHERE assets > ALL (SELECT assets
                    FROM Branch
                    WHERE b-city = "New Territory");
```


The branches in N.T.



solution 2:

```
SELECT X.bname
FROM Branch X
WHERE NOT EXISTS (SELECT *
                  FROM Branch Y
                  WHERE Y.b-city = "New Territory"
                  AND Y.assets >= X.assets);
```

The branches in N.T. with assets greater than branch X



Nested Queries, SOME/ANY

- Find names of all branches that have greater assets than some branch located in Central

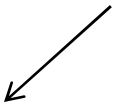
solution 1:

```
SELECT bname  
FROM Branch
```

```
WHERE assets > SOME (SELECT assets  
FROM Branch
```

```
WHERE b-city = "Central");
```

The assets of the
branches in
"Central"



Anyone



solution 2:

```
SELECT X.bname  
FROM Branch X, Branch Y  
WHERE X. assets > Y.assets  
AND Y.b-city= "Central";
```

Aggregate Functions

- Used to summarize information from multiple tuples into a single-tuple
- 1 Grouping (**GROUP BY**)
 - ◆ Create subgroups of tuples before summarizing
 - ◆ E.g., A group of tuples => Count its number to be the value of a new attribute
- 2 Built-in aggregate functions
 - ◆ COUNT, SUM, MAX, MIN, and AVG
 - ◆ Functions can be used in the SELECT clause or in a HAVING clause
- 3 Having :
 - filtering subgroups with conditions

Aggregate Functions

- Examples:
 - assume the following relational schemes:
Borrow (b-name, loan#, c-name, amount)
Customer (c-name, street, city)
Branch (b-name, assets, b-city)
Deposit (b-name, acct#, c-name, balance)

GROUP BY

■ GROUP BY clause

- ◆ Provides a condition to select or reject an entire group:
- ◆ Find the average account balance at each branch

```
SELECT b-name, AVG(balance)
FROM Deposit
GROUP BY b-name;
```

← AVG(balance) is the average of each branch

■ HAVING clause

- ◆ Provides a condition to select or reject an entire group
- ◆ If only interested in branches where average balance is > \$12000

```
SELECT b-name, AVG(balance)
FROM Deposit
GROUP BY b-name
```

CS3402 HAVING AVG(balance) > 12000;

GROUP BY and HAVING Clauses

- **Query 26.** For each project on which more than two employees work, retrieve the project number, the project name, and the number of employees who work on the project.

Q26:	SELECT	Pnumber, Pname, COUNT (*)
	FROM	PROJECT, WORKS_ON
	WHERE	Pnumber=Pno
	GROUP BY	Pnumber
	HAVING	COUNT (*) > 2;

EXPANDED Block Structure of SQL Queries

```
SELECT <attribute and function list>  
FROM <table list>  
[ WHERE <condition> ]  
[ GROUP BY <grouping attribute(s)> ]  
[ HAVING <group condition> ]  
[ ORDER BY <attribute list> ];
```

Views (Virtual Tables)

- CREATE VIEW from SELECT

- ◆ In V1, attributes retain the names from base tables. In V2, attributes are assigned new names

V1: CREATE VIEW WORKS_ON1
 AS SELECT Fname, Lname, Pname, Hours
 FROM EMPLOYEE, PROJECT, WORKS_ON
 WHERE Ssn=Essn AND Pno=Pnumber;

V2: CREATE VIEW DEPT_INFO(Dept_name, No_of_emps, Total_sal)
 AS SELECT Dname, COUNT (*), SUM (Salary)
 FROM DEPARTMENT, EMPLOYEE
 WHERE Dnumber=Dno
 GROUP BY Dname;

↖
New attribute
names

Views

- Example:

```
CREATE VIEW loan-info  
AS SELECT bname, loan#, cname  
FROM Borrow;
```

- Views can be regarded and retrieved as ordinary tables

E.g.,

```
SELECT loan#, cname  
FROM loan-info  
WHERE bname = "Central";
```

Views

- E.g., suppose we define a view:

```
CREATE VIEW branch-city  
AS SELECT bname, city  
FROM Borrow, Customer  
WHERE Borrow.cname = Customer.cname;
```

- Now if we need to insert a tuple (Brighton, Shatin) through this view, it will cause:

(Brighton, null, null, null) => Borrow, and
(null, null, Shatin) => Customer

Then what happens to the following query?

```
SELECT * FROM branch-city;
```

(would the tuple (Brighton, Shatin) be part of the answer?)

Views

- The answer is NO, because all comparisons involving *null* values are defined to be false!
- To avoid such problems and to simplify implementation, most SQL-based DBMSs restrict the following condition:

“A modification is permitted through a view ONLY IF the view is defined in terms of ONE base relation.”

NULL Values

- Information can be very often incomplete in the real world
- Unknown attributes are assigned a **null** value
- One proposal to deal with null values is by using 3-valued logic:

Not	
T	F
U	U
F	T

AND	T	U	F
T	T	U	F
U	U	U	F
F	F	F	F

OR	T	U	F
T	T	T	T
U	T	U	U
F	T	U	F

And

- Only T-T returns T
- And-ing F with anything results with F
- The rest is undefined

Or

- Only F-F returns F
- Or-ing T with anything results with T
- The rest is undefined

NULL Values Comparisons

Condition	Value of a	Evaluation
a IS NULL	10	FALSE
a IS NOT NULL	10	TRUE
a IS NULL	NULL	TRUE
a IS NOT NULL	NULL	FALSE
a = NULL	10	UNKNOWN
a = !NULL	10	UNKNOWN
a = NULL	NULL	UNKNOWN
a = !NULL	NULL	UNKNOWN
a = 10	NULL	UNKNOWN
a = 10	NULL	UNKNOWN

NULL Values

- Any problem with the following query?

```
SELECT c-name  
FROM Deposit  
WHERE balance > 10000  
       OR balance <= 10000;
```

(In fact, all comparisons involving Null become FALSE!)

- In most SQL-based DBMSs, the special keyword NULL may be used to test for a null value. E.g.,

```
SELECT c-name  
FROM Deposit  
WHERE balance IS NULL;  
       (or balance IS NOT NULL;)
```

Data Manipulation Language(DML)

INSERT Command

- Specify the relation name and a list of values for the tuple

```
U1:  INSERT INTO  EMPLOYEE
      VALUES      ( 'Richard', 'K', 'Marini', '653298653', '1962-12-30', '98
                    Oak Forest, Katy, TX', 'M', 37000, '653298653', 4 );
```

```
U3B:  INSERT INTO  WORKS_ON_INFO ( Emp_name, Proj_name,
                                     Hours_per_week )
      SELECT        E.Lname, P.Pname, W.Hours
      FROM          PROJECT P, WORKS_ON W, EMPLOYEE E
      WHERE         P.Pnumber=W.Pno AND W.Essn=E.Ssn;
```

The values for the attributes are obtained from the results of the SELECT statement

DELETE Command

- Removes tuples from a relation

DELETE FROM table-name;

- ◆ (Note: this operation only deletes all tuples from the table and the table is still there)
- ◆ Includes a **WHERE** clause to select the tuples to be deleted

U4A:	DELETE FROM	EMPLOYEE
	WHERE	Lname='Brown';
U4B:	DELETE FROM	EMPLOYEE
	WHERE	Ssn='123456789';
U4C:	DELETE FROM	EMPLOYEE
	WHERE	Dno=5;
U4D:	DELETE FROM	EMPLOYEE;

UPDATE Command

- Modify attribute values of one or more selected tuples
- Additional SET clause in the UPDATE command
 - ◆ Specifies attributes to be modified and new values

```
U5:  UPDATE PROJECT
      SET    Plocation = 'Bellaire', Dnum = 5
      WHERE  Pnumber=10;
```

UPDATE Command

- *Update*

- Example1:

UPDATE Deposit

SET balance = balance * 1.05

(to increase the payment by 5% to all accounts; it is applied to each tuple exactly once.)

- Example2:

UPDATE Deposit

SET balance = balance * 1.06 WHERE balance > 10000

UPDATE Deposit

SET balance = balance * 1.05 WHERE balance <= 10000

(to increase the payment by 6% to all accounts with balance over \$10000; all others receive 5% increase.)

Summary of SQL Syntax

Table 7.2 Summary of SQL Syntax

```
CREATE TABLE <table name> ( <column name> <column type> [ <attribute constraint> ]  
                             { , <column name> <column type> [ <attribute constraint> ] }  
                             [ <table constraint> { , <table constraint> } ] )
```

```
DROP TABLE <table name>  
ALTER TABLE <table name> ADD <column name> <column type>
```

```
SELECT [ DISTINCT ] <attribute list>  
FROM ( <table name> { <alias> } | <joined table> ) { , ( <table name> { <alias> } | <joined table> ) }  
[ WHERE <condition> ]  
[ GROUP BY <grouping attributes> [ HAVING <group selection condition> ] ]  
[ ORDER BY <column name> [ <order> ] { , <column name> [ <order> ] } ]
```

```
<attribute list> ::= ( * | ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) )  
                  { , ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) ) } )
```

```
<grouping attributes> ::= <column name> { , <column name> }
```

```
<order> ::= ( ASC | DESC )
```

```
INSERT INTO <table name> [ ( <column name> { , <column name> } ) ]  
( VALUES ( <constant value> , { <constant value> } ) { , ( <constant value> { , <constant value> } ) }  
| <select statement> )
```

Summary of SQL Syntax

Table 7.2 Summary of SQL Syntax

DELETE FROM <table name>
[WHERE <selection condition>]

UPDATE <table name>
SET <column name> = <value expression> { , <column name> = <value expression> }
[WHERE <selection condition>]

CREATE [UNIQUE] INDEX <index name>
ON <table name> (<column name> [<order>] { , <column name> [<order>] })
[CLUSTER]

DROP INDEX <index name>

CREATE VIEW <view name> [(<column name> { , <column name> })]
AS <select statement>

DROP VIEW <view name>

NOTE: The commands for creating and dropping indexes are not part of standard SQL.

References

- 6e
 - Ch. 4. p. 83 – 107
 - Ch. 5, p. 111 – 126