

---

# ***CS3402 : Chapter 8***

## ***Files and Hashed Files***

# *Storage Medium for Databases*

---


- Memory hierarchy
  - ◆ CPU cache => main memory => flash memory/Phase change memory => magnetic disks/optical disks
  - ◆ Slower in access delay but larger in memory size (less expensive)
  - ◆ Volatile vs. non-volatile
- Primary storage
  - ◆ The storage media that can be operated directly by the CPU
  - ◆ Include main memory and cache memory
- Secondary and tertiary storage
  - ◆ Slower in access
  - ◆ Include magnetic disks, optical disks and flash memory
- A database could be huge in size (several GB or even larger)
  - ◆ Need to be resided in secondary/tertiary memory (**persistent data and non-volatile storage**)

# Disk Storage Devices

---

- Preferred secondary storage device for high storage capacity and low cost
- Data are stored as magnetized areas on magnetic disk surfaces
- A disk pack contains several magnetic disks connected to a rotating spindle
- Disks are divided into concentric circular tracks on each disk surface
  - ◆ Track capacities vary typically from 4 to 50 Kbytes or more
- A track is divided into fixed size sectors and then into blocks
  - ◆ Typical block sizes range from B=512 bytes to B=4096 bytes
  - ◆ Whole blocks are transferred between disk and main memory for processing
- Track -> sectors -> blocks

# Disk Storage Devices




**Figure 17.2**

Different sector organizations on disk. (a) Sectors subtending a fixed angle. (b) Sectors maintaining a uniform recording density.

# Disk Storage Devices

**Figure 17.1**

- (a) A single-sided disk with read/write hardware.
- (b) A disk pack with read/write hardware.

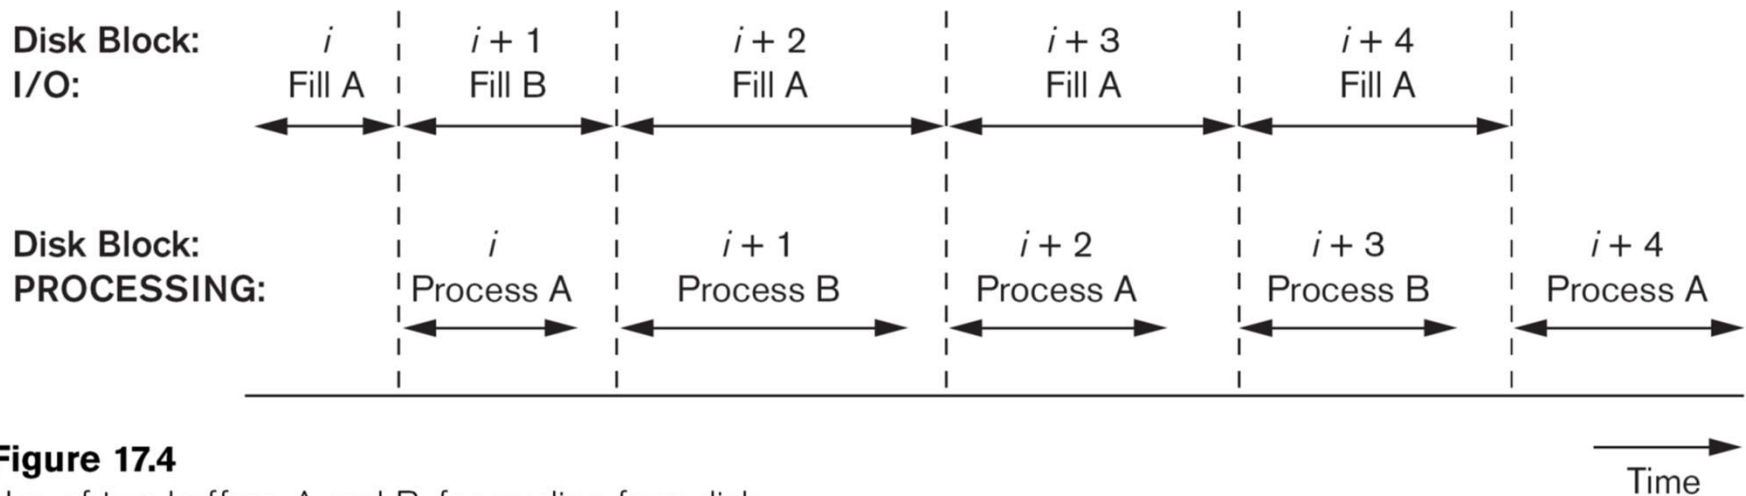


# Disk Storage Devices

---

- A read-write head moves to the track that contains the block to be transferred
  - ◆ Disk rotation moves the block under the read-write head for reading or writing
- To access a physical disk block:
  - ◆ to the identified track number (seek time, e.g., 3 to 8ms)
  - ◆ the block number (within the cylinder) (rotational delay, e.g., 2ms)
  - ◆ and get the block data (transfer delay)
- Disk access delay = seek time + rotational delay + transfer delay
- Reading or writing a disk block is time consuming because of the seek time and rotational delay (latency)
- Double buffering can be used to speed up the transfer of contiguous disk blocks

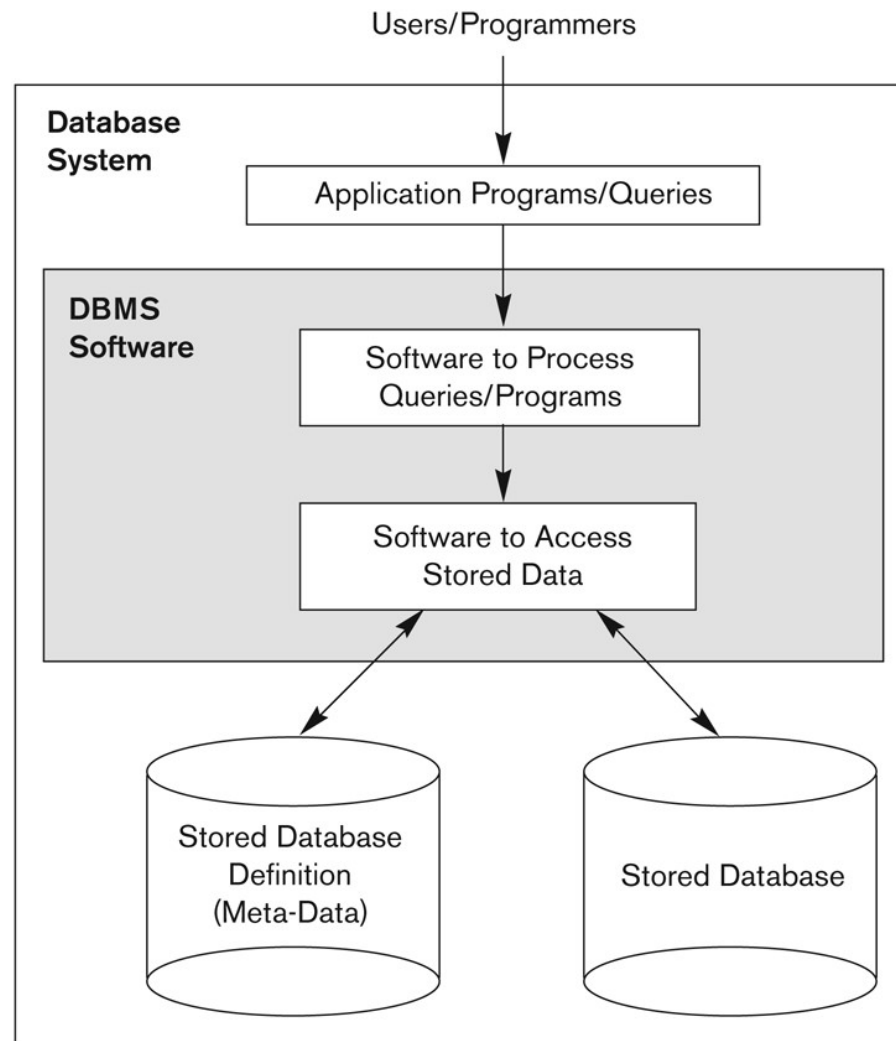
# Double Buffer to Reduce Access Delay



**Figure 17.4**

Use of two buffers, A and B, for reading from disk.

# Database Systems



Store the values  
in files

**Figure 1.1**  
A simplified database  
system environment.




# Records


---

- Database: data file (records) + meta-data
- Fixed and variable length records
- Records contain fields (attributes)
- Fields may be fixed length or variable length (e.g., Varchar)
- Variable length fields can be mixed into one record:
  - ◆ **Separator characters** or length fields are needed so that the record can be “parsed”
- **Blocking:**
  - ◆ Refers to storing a number of records into one block on the disk
- **Blocking factor (bfr)** refers to the number of records per block
- There may be empty space in a block if an integral number of records do not fit into one block

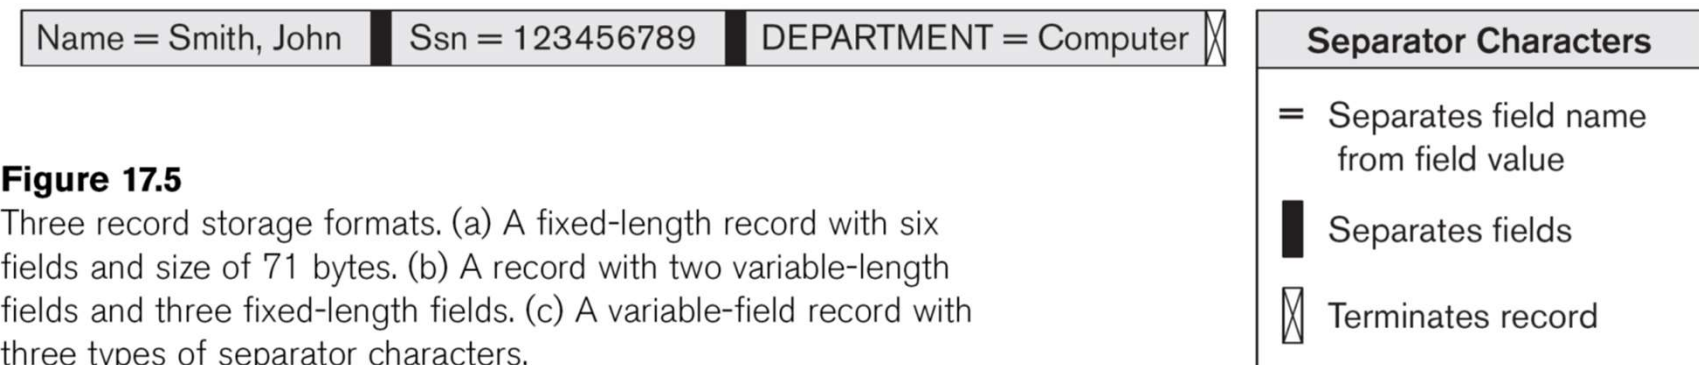
(a)



(b)



(c)



**Figure 17.5**

Three record storage formats. (a) A fixed-length record with six fields and size of 71 bytes. (b) A record with two variable-length fields and three fixed-length fields. (c) A variable-field record with three types of separator characters.

# Files of Records

---

- A file (e.g., table) is a *sequence of records* (e.g., tuples), where each record is a collection of data values (fields)
- A file can have *fixed-length* records or *variable-length* records
- A *file descriptor* (or file header) includes information that describes the file, such as the *field names* and their *data types*, and the *addresses* of the file blocks on disk
- File records are stored on disk blocks

# Files of Records

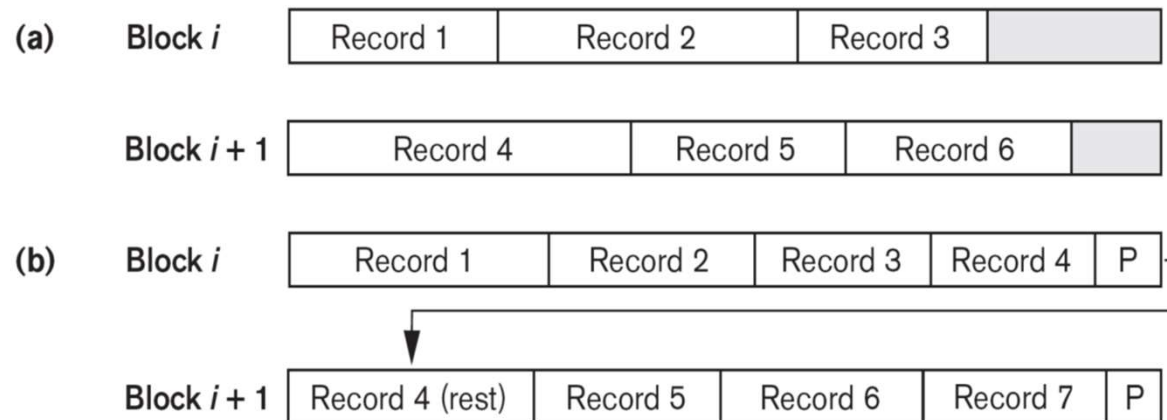
---

- File records can be **unspanned** or **spanned**
  - ◆ Unspanned: no record can span two blocks
  - ◆ Spanned: a record can be stored in more than one block
- The physical disk blocks that are allocated to hold the records of a file can be *contiguous (one by one)*, *linked (using pointers)*, or *indexed (a table to describe their locations)*
- In a file of fixed-length records, all records have the same format. Usually, unspanned blocking is use
- Files of variable-length records require additional information to be stored in each record, such as **separator characters**
  - ◆ Usually spanned blocking is used with such files

# Spanned Vs. Unspanned Files of Records

**Figure 17.6**

Types of record organization.  
(a) Unspanned.  
(b) Spanned.



Suppose the block size is  $B$ . For fixed-length records of size  $R$  with  $B \geq R$ , blocking factor  $(bfr) = B/R$  round down  
The unused space in each block =  $B - (bfr \times R)$  bytes

# *Typical Operations on Files*

---

- Typical operations on files:
  - ◆ **OPEN**: makes the file ready for access, and associates a **pointer** that will refer to a *current* file record at each point in time
  - ◆ **FIND**: searches for the first file record that satisfies a certain condition, and makes it the current file record
  - ◆ **FINDNEXT**: searches for the next file record (from the current record) that satisfies a certain condition, and makes it the current file record
  - ◆ **READ**: reads the current file record into a **program variable**
  - ◆ **INSERT**: inserts a new record into the file, and makes it the current file record

# *Typical Operations on Files*

---

- Typical operations on files:
  - ◆ **DELETE**: removes the current file record from the file, usually by **marking** the record to indicate that it is no longer valid
  - ◆ **MODIFY**: changes the values of some fields of the current file record
  - ◆ **CLOSE**: terminates access to the file
  - ◆ **REORGANIZE**: reorganizes the file records. For example, the records marked “deleted” are physically removed from the file or a new organization of the file records is created
  - ◆ **READ\_ORDERED**: reads the file blocks in order of a specific field of the file

# Unordered Files

---

- Also called a **heap** file (records are unordered)
- New records are **inserted at the end** of the file
  - ◆ Arranged in their insertion sequence
- A **linear search** through the file records is necessary to search for a record
  - ◆ This requires reading and searching half the file blocks on **average**, and is hence quite expensive ( $n/2$ )
  - ◆ Worst case, all records ( $n$ )
- Record insertion is efficient (add to the end)
- Reading the records in order of a particular field requires sorting the file records

9	16	50	2	10	4	8	12	60	100
---	----	----	---	----	---	---	----	----	-----



# Ordered Files

---

- Also called a **sequential** file (records are ordered)
- File records are kept sorted by the values of an **ordering field**
- **Insertion is expensive**: records must be inserted in the correct order
  - ◆ It is common to keep a separate **unordered overflow** file for new records to improve insertion efficiency; this is periodically merged with the main ordered file
- A **binary search** can be used to search for a record on its **ordering field** value
  - ◆ This requires reading and searching  $\log_2 n$  of the file blocks on the average, an improvement over linear search
- Reading the records in **order of the ordering field** is quite **efficient**

2	4	8	9	10	12	16	50	60	100
---	---	---	---	----	----	----	----	----	-----

# Ordered Files


---

Ordered file

2	4	8	9	10	12	16	50	60	100
---	---	---	---	----	----	----	----	----	-----




merging



18	20	11	46	71
----	----	----	----	----

Unordered overflow  
buffer

2	4	8	9	10	11	12	20	28	46
---	---	---	---	----	----	----	----	----	----



# Ordered Files

	NAME	SSN	BIRTHDATE	JOB	SALARY	SEX
block 1	Aaron, Ed					
	Abbott, Diane					
	⋮					
	Acosta, Marc					
block 2	Adams, John					
	Adams, Robin					
	⋮					
	Akers, Jan					
block 3	Alexander, Ed					
	Alfred, Bob					
	⋮					
	Allen, Sam					
block 4	Allen, Troy					
	Anders, Keith					
	⋮					
	Anderson, Rob					
block 5	Anderson, Zach					
	Angeli, Joe					
	⋮					
	Archer, Sue					
block 6	Arnold, Mack					
	Arnold, Steven					
	⋮					
	Atkins, Timothy					
⋮						
block n - 1	Wong, James					
	Wood, Donald					
	⋮					
	Woods, Manny					
block n	Wright, Pam					
	Wyatt, Charles					
	⋮					
	Zimmer, Byron					

# Binary Search

---

**Algorithm 17.1.** Binary Search on an Ordering Key of a Disk File

$l \leftarrow 1; u \leftarrow b$ ; (\*  $b$  is the number of file blocks \*)

while ( $u \geq l$ ) do

**begin**  $i \leftarrow (l + u) \text{ div } 2$ ;

    read block  $i$  of the file into the buffer;

    if  $K < (\text{ordering key field value of the } \textit{first} \text{ record in block } i)$

        then  $u \leftarrow i - 1$

    else if  $K > (\text{ordering key field value of the } \textit{last} \text{ record in block } i)$

        then  $l \leftarrow i + 1$

    else if the record with ordering key field value =  $K$  is in the buffer

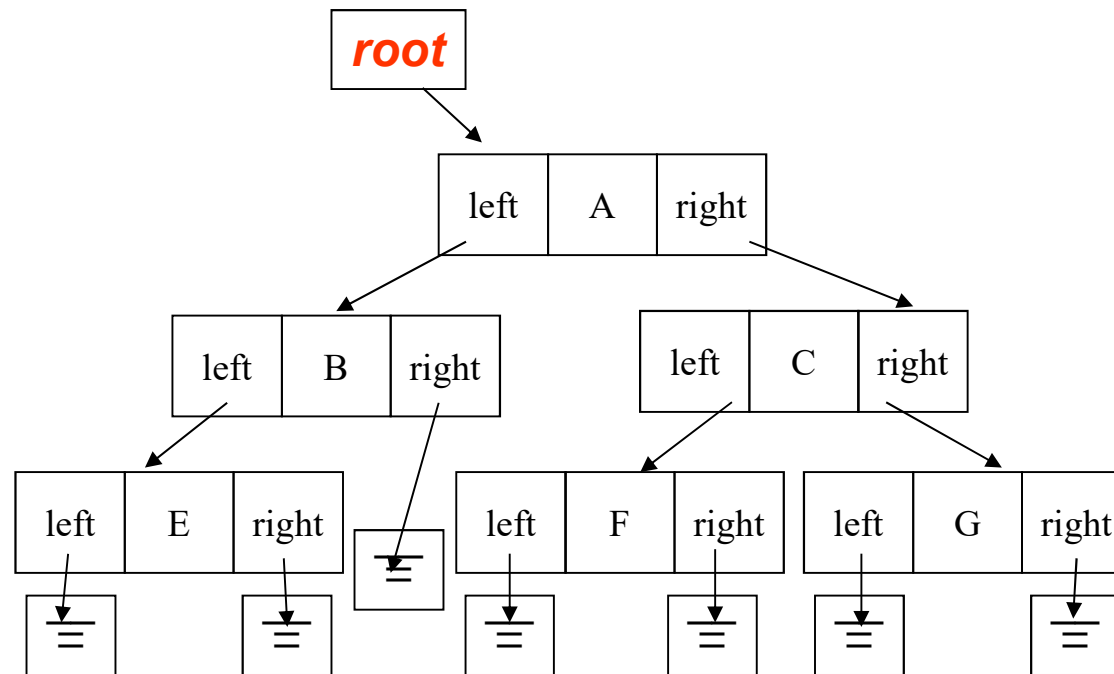
        then goto found

    else goto notfound;

**end**;

goto notfound;

# Binary Tree Data Structures



```
class TreeNode
{
private:
    int info;
    TreeNode* left;
    TreeNode* right;
};
class Mytree
{
private:
    TreeNode* root;
}
```

- **Binary Tree:** every node has two links (pointers)

# Average Access Times

---

- The following table shows the average access time to access a specific record for a given type of file

**Table 17.2** Average Access Times for a File of  $b$  Blocks under Basic File Organizations

Type of Organization	Access/Search Method	Average Blocks to Access a Specific Record
Heap (unordered)	Sequential scan (linear search)	$b/2$
Ordered	Sequential scan	$b/2$
Ordered	Binary search	$\log_2 b$


# Hashed Files

---

- Hashing for disk files is called External Hashing (files on disk)
- The file blocks are divided into  $M$  equal-sized buckets, numbered  $\text{bucket}_0, \text{bucket}_1, \dots, \text{bucket}_{M-1}$
- One of the file fields is designated to be the hash key of the file
- Suppose there is a hash function that takes a hash key as an argument to compute an integer in the range  $0$  to  $B - 1$  where  $B$  is the number of buckets
- A bucket array (an array index) from  $0$  to  $B - 1$  holds the headers of  $B$  lists, one for each bucket of the array
- If a record has search key  $K$ , we store the record by linking it to the bucket list for the bucket numbered  $h(K)$  where  $h$  is the hash function
  - ◆  $h(K) \Rightarrow 0 \text{ to } B - 1$


# Hashed Files

---





# External Hashed Files




**Figure 17.9**

Matching bucket numbers to disk block addresses.

# Example Hashed Tables

- We assume a block can hold two records and  $B = 4$ , i.e., the hash function  $h$  returns values from 0 to 3
- Suppose we add to the hash table a record with key  $g$  and  $h(g) = 1$ . Then we add the new record to the bucket numbered 1
- **Collisions** occur when a new record hashes to a bucket that is already full




# Hashed Files – Collision Resolution


---

- There are numerous methods for collision resolution:
  - ◆ **Open addressing**: proceeding from the occupied position specified by the hash address, the program checks the subsequent positions in order until an unused (empty) position is found.
    - ◆ Linear Probing
      - ◆ If collide, try Bucket\_id+1, Bucket\_id+2, ..., Bucket\_id+n
    - ◆ Quadratic Probing
      - ◆ If collide, try Bucket\_id+1, Bucket\_id+4, ..., Bucket\_id+n<sup>2</sup>
  - ◆ What are the differences in performance?

# Hashed Files – Collision Resolution



Linear Probing



Quadratic Probing

Assume  $h(x)=x\%10$  (take the last digit), and every slot is a bucket

# Hashed Files Collision Resolution

---

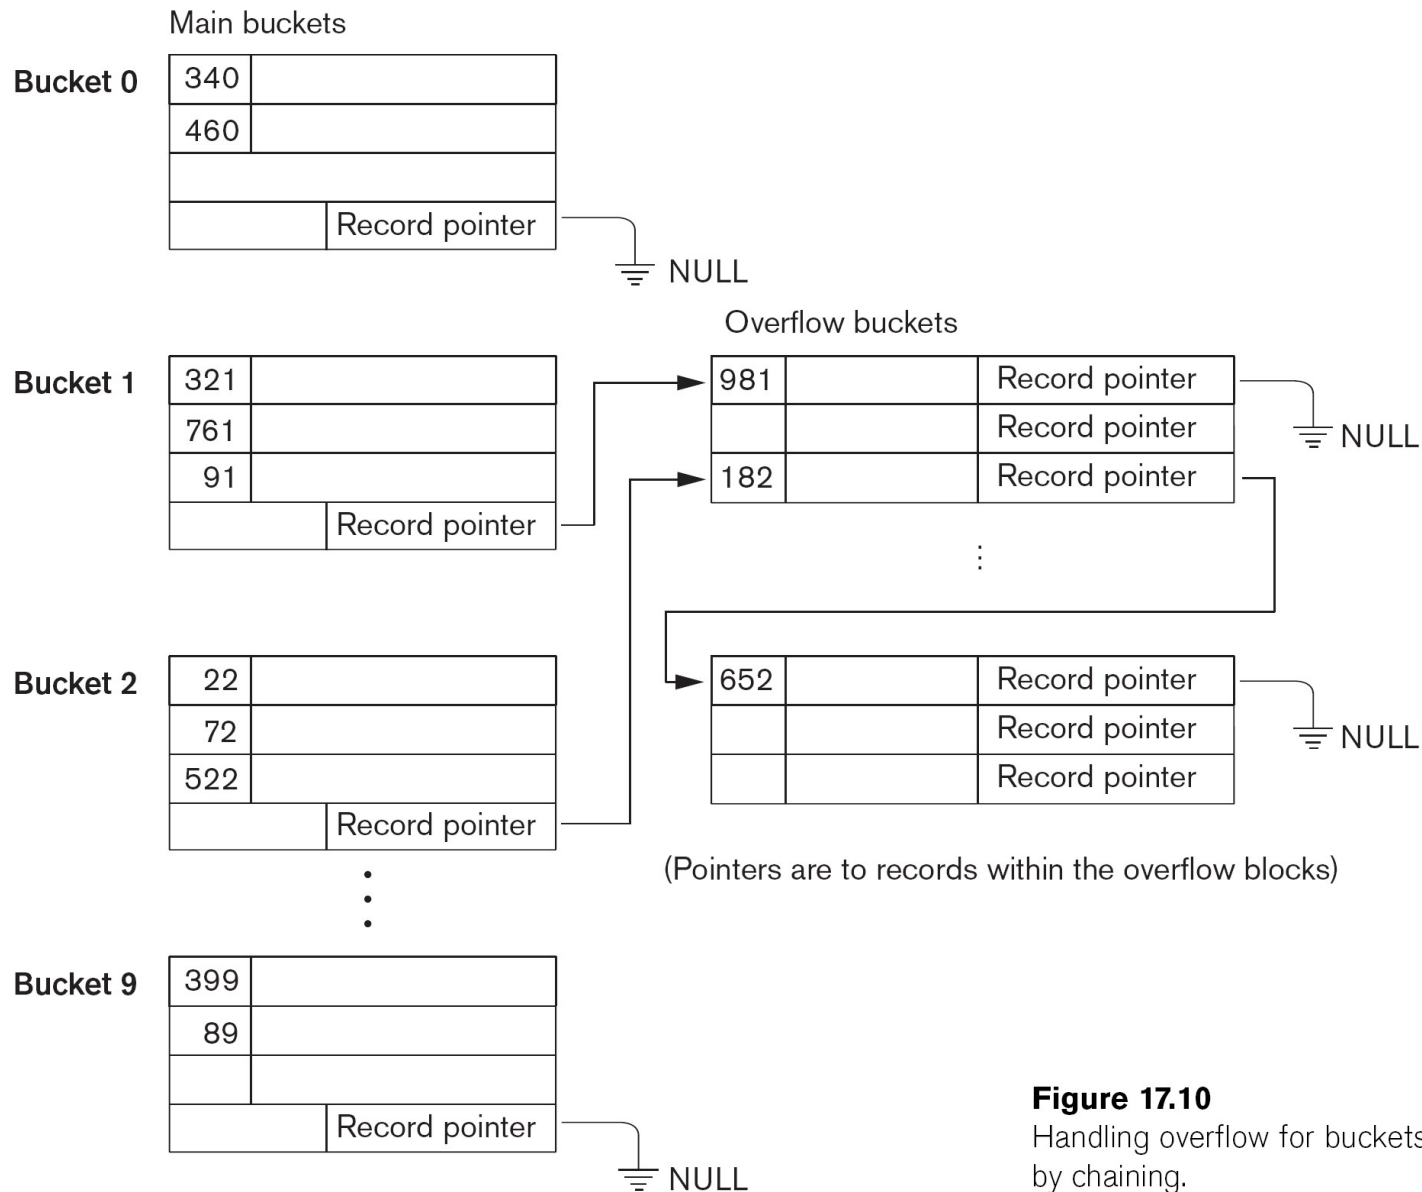
## ■ Chaining:

- ◆ For this method, various overflow locations are kept, usually by extending the array with a number of overflow positions
- ◆ In addition, a **pointer field** is added to each bucket
- ◆ A collision is resolved by placing the new record in an unused overflow bucket and setting the pointer of the occupied hash address bucket to the address of that overflow bucket

## ■ Multiple hashing:

- ◆ The program applies a second hash function if the first results in a collision
- ◆ If another collision results, the program uses open addressing or applies a third hash function and then uses open addressing if necessary

# Hashed Files - Overflow Handling



**Figure 17.10**  
Handling overflow for buckets  
by chaining.

# Hashed Files

---

- To reduce the number of overflow, a hash file is typically kept 70-80% full
- The hash function  $h$  should distribute the records uniformly among the buckets
  - ◆ Otherwise, search time will be increased because many overflow records will exist
  - ◆ Searching overflow records are more expensive
- Main disadvantages of static external hashing:
  - ◆ Fixed number of buckets  $M$  is a problem if the number of records in the file grows or shrinks
  - ◆ Ordered access on the hash key is quite inefficient (requires sorting the records)

# *Extendible and Dynamic Hashing*

---

- Dynamic and Extendible Hashing Techniques
  - ◆ Hashing techniques are extended to allow dynamic growth and shrinking of the number of file records
  - ◆ These techniques include the following: dynamic hashing and extendible hashing
- Both dynamic and extendible hashing use the binary representation (e.g., 1100...) of the hash value  $h(K)$  in order to access a directory
  - ◆ In dynamic hashing the directory is a binary tree
  - ◆ In extendible hashing the directory is an array of size  $2^d$  where  $d$  is called the global depth



# *Extendible and Dynamic Hashing*

---

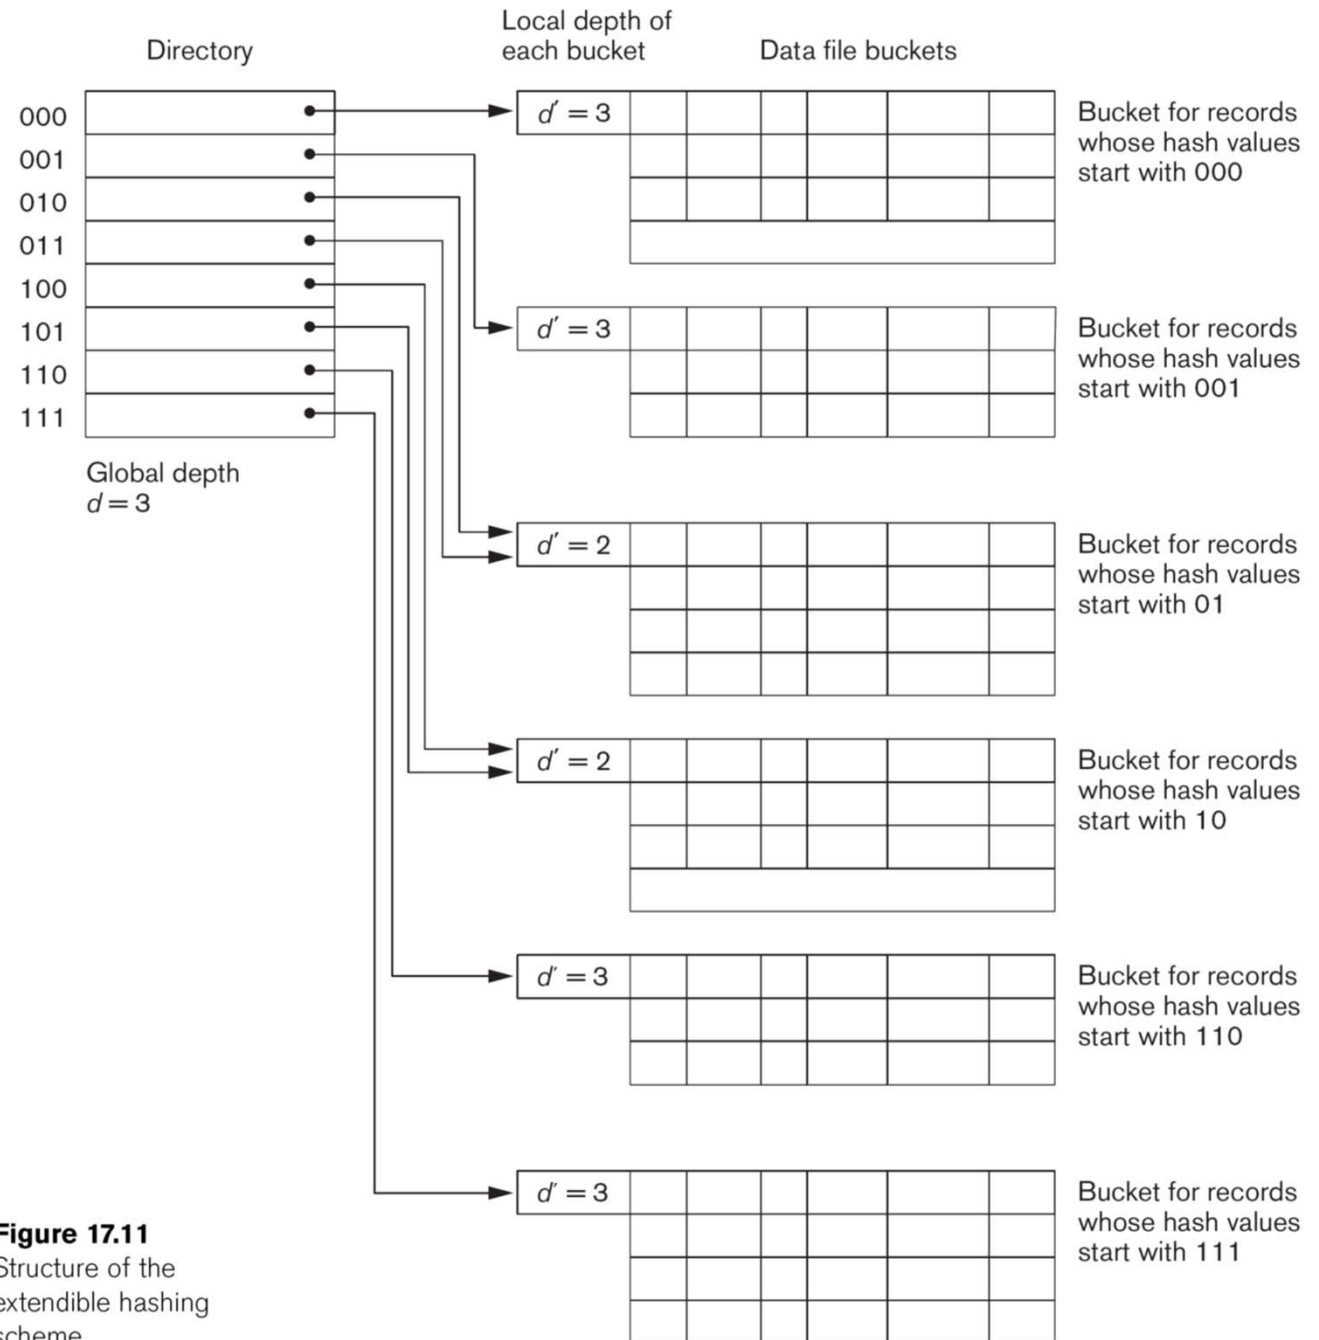
- The **directories** can be stored on disk, and they expand or shrink dynamically
  - ◆ **Directory entries** point to the disk blocks that contain the stored records
- An insertion in a disk block that is full causes the block to **split** into two blocks and the records are redistributed among the two blocks
  - ◆ The directory is updated appropriately

# Extendible Hashing

---

- A directory consisting of an array of  $2^d$  bucket addresses is maintained
- $d$  is called the **global depth** of the directory
- The integer value corresponding to the **first (high-order)  $d$**  bits of a hash value is used as an index to the array to determine a directory entry and the address in that entry determines the bucket storing the records
- A **location  $d'$**  (called, local depth stored with each bucket) specifies the number of bits on which the bucket contents are based
- The value of  $d'$  can be increased or decreased by one at a time to handle overflow or underflow respectively

# Extendible Hashing



**Figure 17.11**  
Structure of the  
extendible hashing  
scheme.


# Extendible Hashing Example

---

- Suppose  $d = 4$ , i.e., the hash function produces a sequence of four bits
- At the moment, only one of these bits is used as illustrated by  $i = 1$  ( $d' = 1$ ) in the box above the bucket array
- The bucket array therefore has only two entries and points to two blocks
  - ◆ The first holds all the current records whose search keys hash to a bit sequence **beginning with 0**
  - ◆ The second holds all those whose search keys hash to a sequence **beginning with 1**
- Suppose we insert a record whose key hash to the sequence 1010
- Since the first bit is 1, it belongs in the second block
- However, the second block is full. It needs to be split by setting  $d'$  to 2

# Extendible Hashing Example

- The two entries beginning with 0 each point to the block for records whose hashed keys begin with 0 and the block still has the integer 1 in its “nub” to indicate that only the first bit determine membership in the block
- The blocks for records beginning with 1 needs to be split into 10 and 11




# Dynamic Hashing

---

- Dynamic and extendible hashing **do not require** an overflow area in general
- Dynamic hashing maintains tree-structured directory with two types of nodes
  - ◆ **Internal nodes** that have two pointers: the left pointer corresponding to the 0 bit (in the hash address) and a right pointer corresponding to the 1 bit
  - ◆ **Leaf nodes**: these hold a pointer to the actual bucket with records

# Dynamic Hashing



**Figure 17.12**  
Structure of the dynamic hashing scheme.

# Reference

---

- 6e
  - ◆ *Ch. 16, pages 565-598*