

AST20105 Data Structures & Algorithms

CHAPTER 7 –TREES II

Instructed by Garret Lai

Before Start


- ▶ At the previous chapter, we learnt about
 - ▶ Binary Search Tree
- ▶ Indeed, trees can have more than two children.
This tree is called
 - ▶ A multiway tree of order m ,
or
 - ▶ An m -way tree.

Introduction

- ▶ Binary Search Trees provide $O(\log n)$ searching provided that the nodes are being placed in “balanced” manner
- ▶ However, this is not always the case as insertion and deletion of tree nodes will generally make the resulting trees unbalanced
- ▶ In the worst case, the tree is de-generated to a sorted linked list and the searching time is $O(n)$
- ▶ Target:
 - ▶ A Binary Search Tree with n node has height $O(\log n)$

AVL (Adelson-Velsky and Landis) Trees

- ▶ An **AVL tree** is a **Binary Search Tree** where the **height** of the two sub-trees of **ANY** node differs by at most one
- ▶ Each node stores a height value, which refers the height of the node that will be used for balancing check



Every sub-tree of an AVL tree is an AVL tree

```
class AVLNode
{
    public:
        double data;
        int height;
        AVLNode* left;
        AVLNode* right;
};
```

AVLNode.h

AVL Trees

- ▶ With this property, AVL trees are balanced and it is guaranteed that **height is logarithmic** in the number of items (denoted as n) in the trees, i.e. $\log(n)$
- ▶ Therefore, efficiency of the tree operations can always be guaranteed
 - ▶ Searching: $O(\log n)$ in the worst case
 - ▶ Insertion: $O(\log n)$ in the worst case
 - ▶ Deletion: $O(\log n)$ in the worst case

AVL Tree – Searching

- Searching in AVL Trees is the **same** as the one for Binary Search Tree introduced in the last topic


AVL Tree – Insertion

- ▶ To insert an element in an AVL tree
 - ▶ **Search** the tree using a search algorithm similar to the one for Binary Search Tree and find the place where the new item should be inserted
 - ▶ **Create a new node** with the element and **attached** it to the tree
- ▶ The insertion may cause the AVL tree **imbalance**, we must perform a **tree “rotation”** to fix the imbalance
- ▶ Types of rotation
 - ▶ **Single** rotation
 - ▶ **Double** rotation
(i.e. two single rotations)


AVL Tree – Insertion

- ▶ When a **node is un-balanced**,
4 cases need to be considered
 - ▶ Left-left: The insertion was in the left sub-tree of the left child of the un-balanced node
 - ▶ Single rotation is required
 - ▶ Right-right: The insertion was in the right sub-tree of the right child of the un-balanced node
 - ▶ Single rotation is required
 - ▶ Left-right: The insertion was in the right sub-tree of the left child of the un-balanced node
 - ▶ Double rotation is required
 - ▶ Right-left: The insertion was in the left sub-tree of the right child of the un-balanced node
 - ▶ Double rotation is required

AVL Tree – Left-left Single Rotation




AVL Tree – Right-right Single Rotation




Left rotation at a

AVL Tree – Left-right Double Rotation




First rotate left at a and
then right at c


First: Left rotation at a



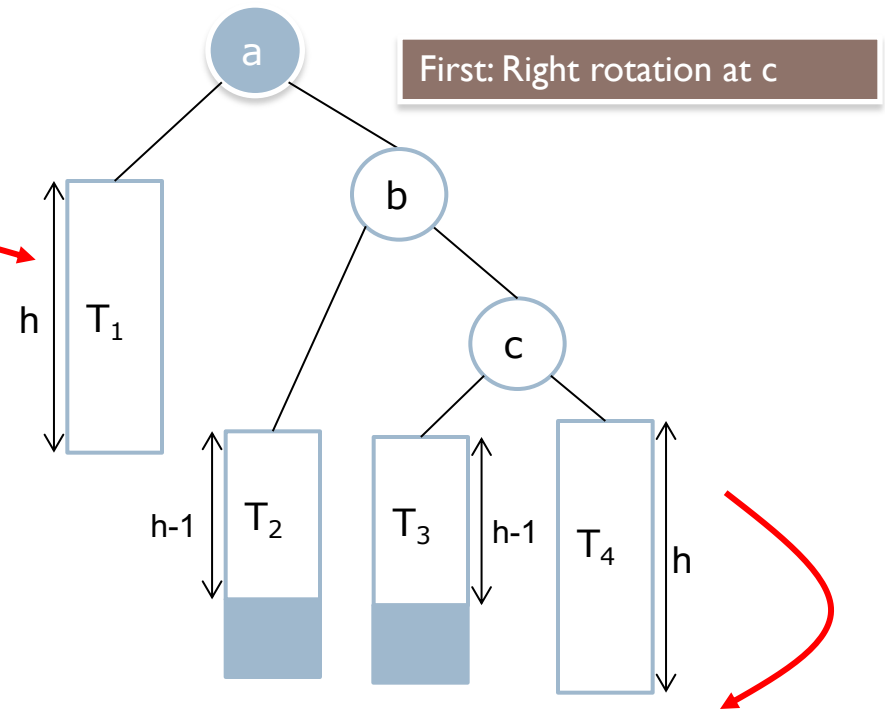
Second: Right rotation at c




AVL Tree – Right-left Double Rotation



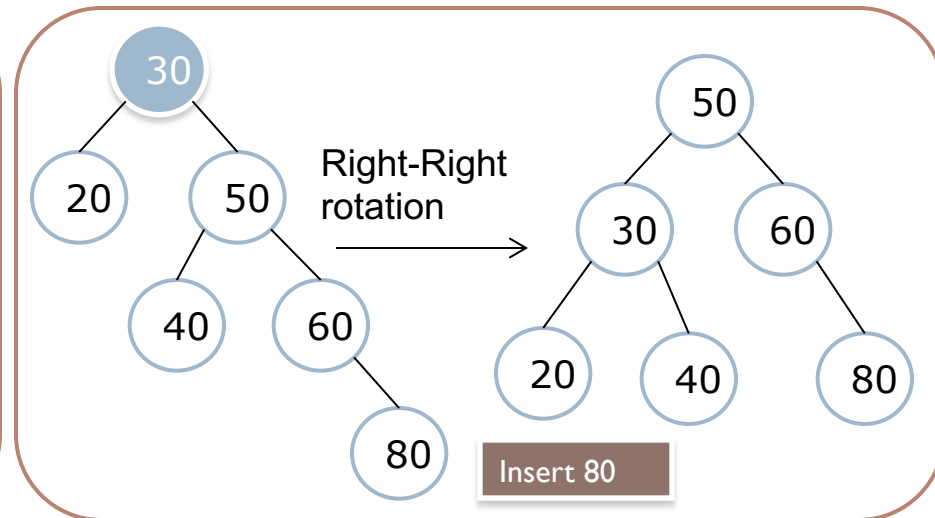
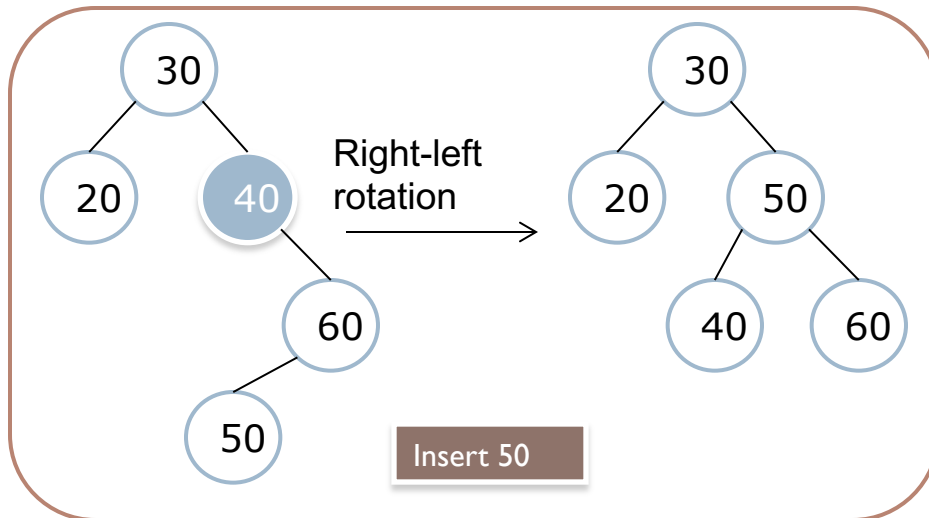




First rotate right at c and
then left rotate at a



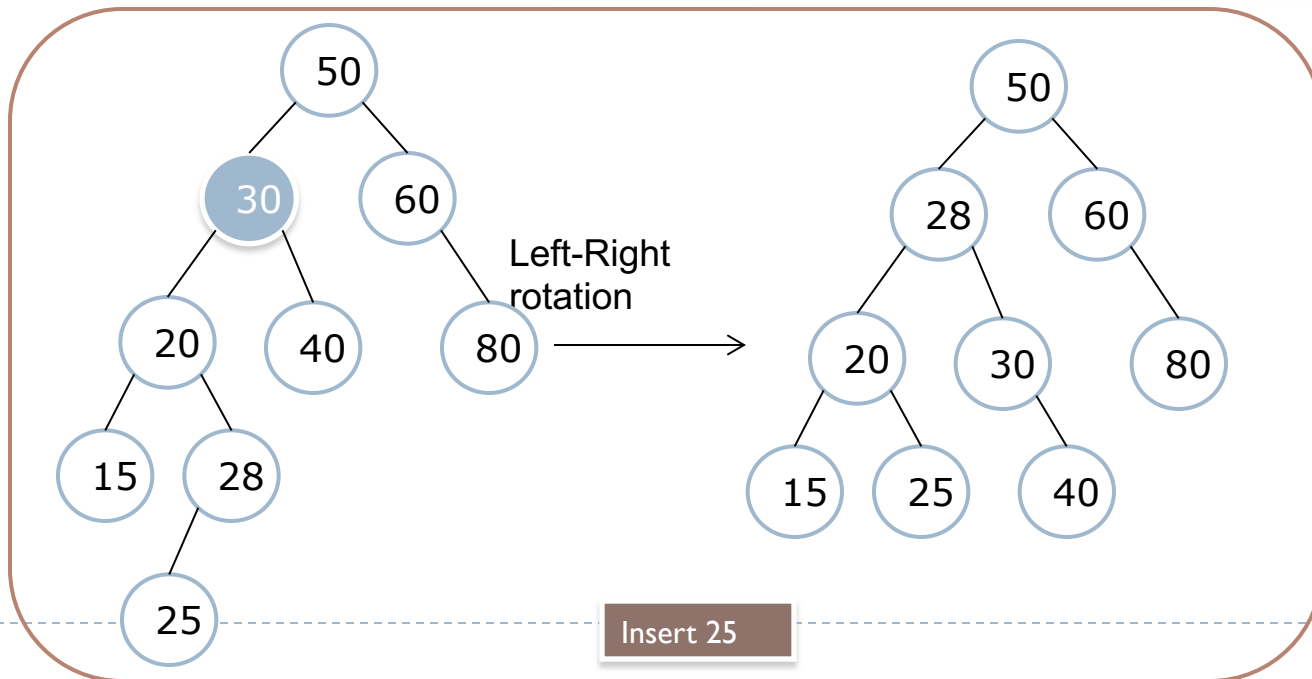



Second: Left rotation at a



Examples



Examples



AVL Tree – C++ Code

```
int AVLTree::height(AVLNode* p)
{
    if(p == NULL)
        return -1;
    else
        return p->height;
}

int AVLTree::max(int l, int r)
{
    if(l > r)
        return l;
    else
        return r;
}
```

AVLNode.cpp

```

AVLNode* AVLTree::singleRotateWithLeft(AVLNode* k2) {
    AVLNode* k1;
    k1 = k2->left;
    k2->left = k1->right;
    k1->right = k2;
    k2->height = max(height(k2->left), height(k2->right)) + 1;
    k1->height = max(height(k1->left), k2->height) + 1;
    return k1;
}

```

```

AVLNode* AVLTree::singleRotateWithRight(AVLNode* k1) {
    AVLNode* k2;
    k2 = k1->right;
    k1->right = k2->left;
    k2->left = k1;
    k1->height = max(height(k1->left), height(k1->right)) + 1;
    k2->height = max(height(k2->right), k1->height) + 1;
    return k2;
}

```

```

AVLNode* AVLTree::doubleRotateWithLeft(AVLNode* k) {
    k->left = singleRotateWithRight(k->left);
    return singleRotateWithLeft(k);
}

```

```

AVLNode* AVLTree::doubleRotateWithRight(AVLNode* k) {
    k->right = singleRotateWithLeft(k->right);
    return singleRotateWithRight(k);
}

```



```

AVLNode* AVLTree::insert(double x, AVLNode* t) {
    if(t == NULL) {
        t = new AVLNode;
        if(t == NULL) {
            cout << "Out of memory" << endl; exit(1);
        }
        else { t->data = x; t->height = 0; t->left = t->right = NULL; }
    }
    else
        if(x < t->data) {
            t->left = insert(x, t->left);
            if(height(t->left) - height(t->right) == 2)
                if(x < t->left->data)
                    t = singleRotateWithLeft(t);
                else
                    t = doubleRotateWithLeft(t);
        }
        else
            if(x > t->data) {
                t->right = insert(x, t->right);
                if(height(t->right) - height(t->left) == 2)
                    if(x > t->right->data)
                        t = singleRotateWithRight(t);
                    else
                        t = doubleRotateWithRight(t);
            }
    t->height = max(height(t->left), height(t->right)) + 1;
    return t;
}

```

AVL Tree – Deletion

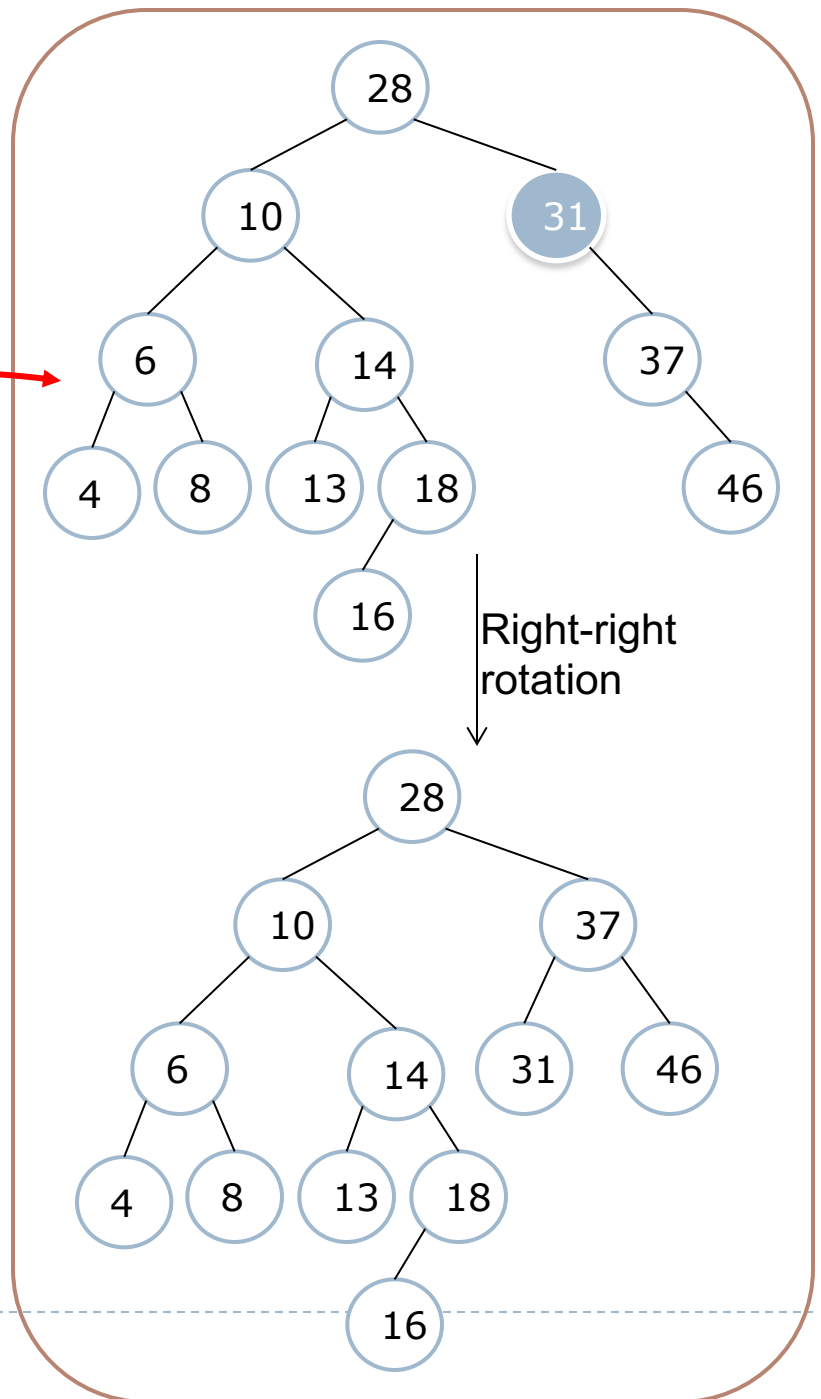

- ▶ To delete an element from an AVL tree
 - ▶ **Search** the tree using a search algorithm similar to the one for Binary Search Tree and find the node with the key should be deleted
- ▶ The deletion may cause the AVL tree imbalance, we must perform a tree “**rotation**” to fix the imbalance
- ▶ Types of rotation
 - ▶ **Single** rotation
 - ▶ **Double** rotation (i.e. two single rotations)

AVL Tree – Deletion


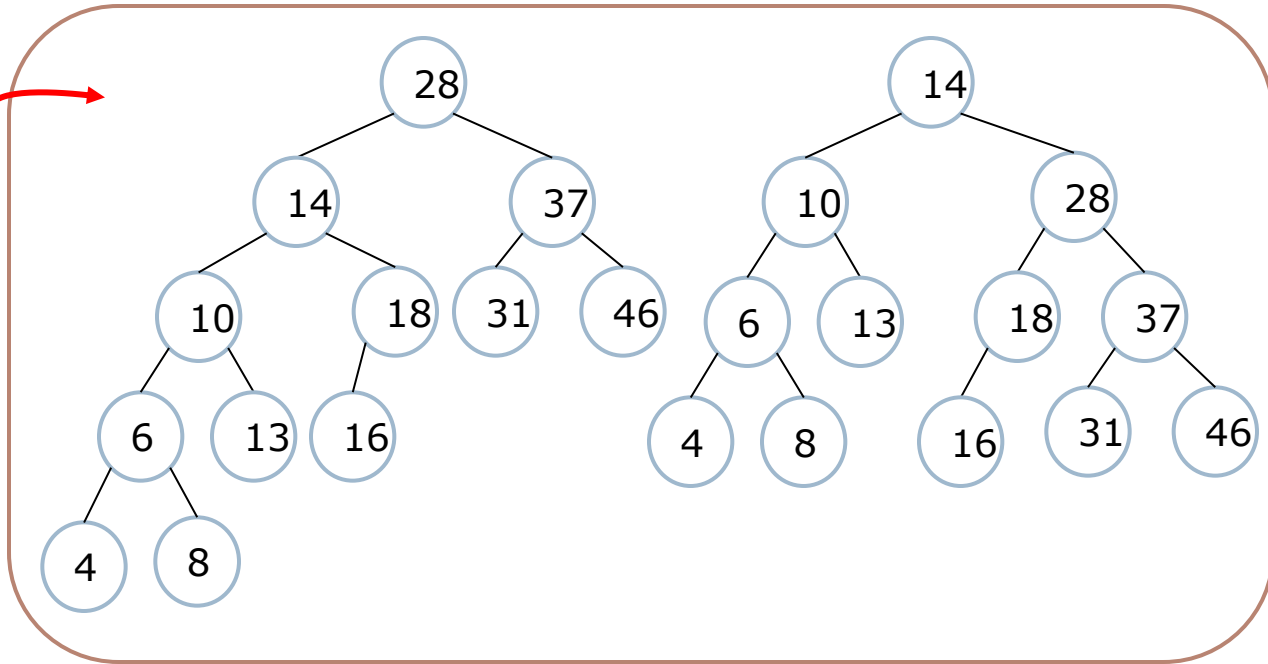

- ▶ When a **node is un-balanced**, **3 cases** need to be considered
 - ▶ The node to be removed with **no children**
 - ▶ **Replace** it with **NULL**
 - ▶ The node to be removed with **one child**
 - ▶ **Replace** it with the **only child**
 - ▶ The node to be removed with **two children**
 - ▶ **Replace** with the **leftmost node in the right sub-tree**
- ▶ Removing a node can **make multiple ancestors unbalanced**, so you are required to **go all the way up to the root** for re-balancing

Examples


Replace with the
leftmost node in
the right sub-tree



Examples



Delete 8



Delete 18



AVL Trees – Pros and Cons

▶ Pros:

- ▶ Time complexity for searching is $O(\log n)$ since AVL trees are always balanced
- ▶ Insertion and deletions are also $O(\log n)$ (Dominated by searching step)
- ▶ The height balancing adds no more than a constant factor to the speed of insertion and deletion

▶ Cons:

- ▶ Difficult to program
- ▶ More space for keeping height of node
- ▶ Intensive searching are done in data systems uses other data structures, e.g. B-trees

Motivation of using B+ Tree

- ▶ AVL Tree is an excellent data structure when the **entire tree can fit into the main memory**
- ▶ When the **data size is large** that the AVL tree cannot fit into the main memory, the performance of **AVL tree operations deteriorated** rapidly. Since pointer operations require disk access, which is generally a bottleneck
- ▶ Therefore, it is important to minimize the number of disk accesses by performing more processor instructions, as processor nowadays is getting faster
- ▶ Idea: Allow a node in a tree **having many children** → B+ tree

Motivation of using B+ Tree

- ▶ The basic unit of I/O operations associated with disk is a **block**.
- ▶ When information is read from a disk, the **entire block** containing this information is **read into memory**,
- ▶ and when information is stored on a disk, an entire block is **written to the disk**.

Motivation of using B+ Tree

- ▶ Each time information is requested from a disk,
 - ▶ this information has to be **located** on the disk,
 - ▶ the **head** has to be **positioned above** the part of the disk where the information resides, and
 - ▶ the disk has to be **spun** so that the entire block passes underneath the head to be **transferred to memory**.

Motivation of using B+ Tree

- ▶ This means that there are several time components for data access:
 - ▶ access time = seek time +
rotational delay (latency) +
transfer time
- ▶ This process is **extremely slow** compared to transferring information within memory.

Motivation of using B+ Tree

- ▶ The first component, **seek time**, is particularly slow because it depends on the **mechanical movement of the disk** head to position the head at the correct track of the disk.
- ▶ **Latency** is the time required to **position the head** above the correct block, and on the average, it is equal to the time needed to make one-half of a revolution.

Motivation of using B+ Tree

- ▶ Transferring information to and from the **disk** is on the order of **milliseconds**.
- ▶ The **CPU processes data** on the order of **microseconds**, 1000 times faster, or on the order of **nanoseconds**, 1 million times faster, or even faster.
- ▶ We can see that processing information on **secondary storage** can significantly **decrease the speed of a program**.

Motivation of using B+ Tree

- ▶ If a program **constantly uses** information stored in secondary storage,
- ▶ the characteristics of this storage have to be taken into account when **designing the program**.

Motivation of using B+ Tree

- ▶ For example, a **binary search tree** can be **spread over** many different **blocks** on a disk.
- ▶ When the tree is **used frequently** in a program, these accesses can **significantly slow down** the execution time of the program.
- ▶ Also, **inserting and deleting** keys in this tree require **many block accesses**.

Motivation of using B+ Tree

- ▶ The BST, which is such an **efficient tool** when it resides entirely **in memory**, turns out to be an encumbrance.


Motivation of using B+ Tree

- ▶ It is better to access a **large amount of data at one time** than to jump from one position on the disk to another to transfer small portions of data.
- ▶ The reason is that each **disk access is very costly**; if possible, the data should be organized to minimize the number of accesses.

Multiway Search Trees

- ▶ A **multiway search tree of order m** , or an m -way search tree, is a multiway tree in which
 - ▶ Each node has **m children** and **$m - 1$ keys**.
 - ▶ The keys in each node are in **ascending order**
 - ▶ The keys in the first i children are **smaller than** the i^{th} key
 - ▶ The keys in the last $m - i$ children are **larger than** the i^{th} key

Multiway Search Trees



Multiway Search Trees

- ▶ The m-way search trees **play the same role** among m-way trees that binary search trees play among binary trees, and
- ▶ They are used for the **same purpose**:
 - ▶ **Fast information retrieval and update.**

B-trees

- ▶ In database programs where most information is stored on disks or tapes, the time penalty for accessing secondary storage can be significantly reduced by the **proper choice of data structures**.
- ▶ **B- trees** (Bayer and McCreight, 1972) are one such approach.

B-trees

- ▶ A B-tree operates closely with secondary storage and can be tuned to **reduce the impediments** imposed by this storage.
- ▶ One important property of B-trees is
 - ▶ The **size of each node**, which can be made as large as the size of a block.
 - ▶ The **number of keys** in one node can vary depending on the **sizes of the keys, organization of the data**, and of course, on the **size of a block**.

B-trees

- ▶ Block size varies for each system.
 - ▶ 512 bytes
 - ▶ 4KB
 - ▶ Or more.
- ▶ Block size is the size of each node of a B-tree.


B-trees

- ▶ A B-tree of order m is a multiway search tree with the following properties:
 - ▶ The root has **at least two subtrees** unless it is a leaf.
 - ▶ Each nonroot and each nonleaf node **holds $k - 1$ keys** and **k pointers** to subtrees where **$\lceil m/2 \rceil \leq k \leq m$** .
 - ▶ Each leaf node holds $k - 1$ keys where $\lceil m/2 \rceil \leq k \leq m$.
 - ▶ All leaves are on the **same level**.


B-trees

- ▶ According to the conditions, a B-tree is always
 - ▶ At least **half full**,
 - ▶ Has **few levels**, and
 - ▶ Is **perfectly balanced**.
 - ▶ A tree is height-balanced or simply balanced if the **difference in height** of both subtrees of any node in the tree is either **zero or one**.

B-trees



Balance



Unbalance

B-trees

- ▶ A node of a B-tree is usually implemented as a **class** containing
 - ▶ an **array of $m - 1$ cells** for keys,
 - ▶ an **m -cell array of pointers** to other nodes, and
 - ▶ possibly **other information** facilitating tree maintenance,
 - ▶ such as the number of keys in a node and a leaf/nonleaf flag

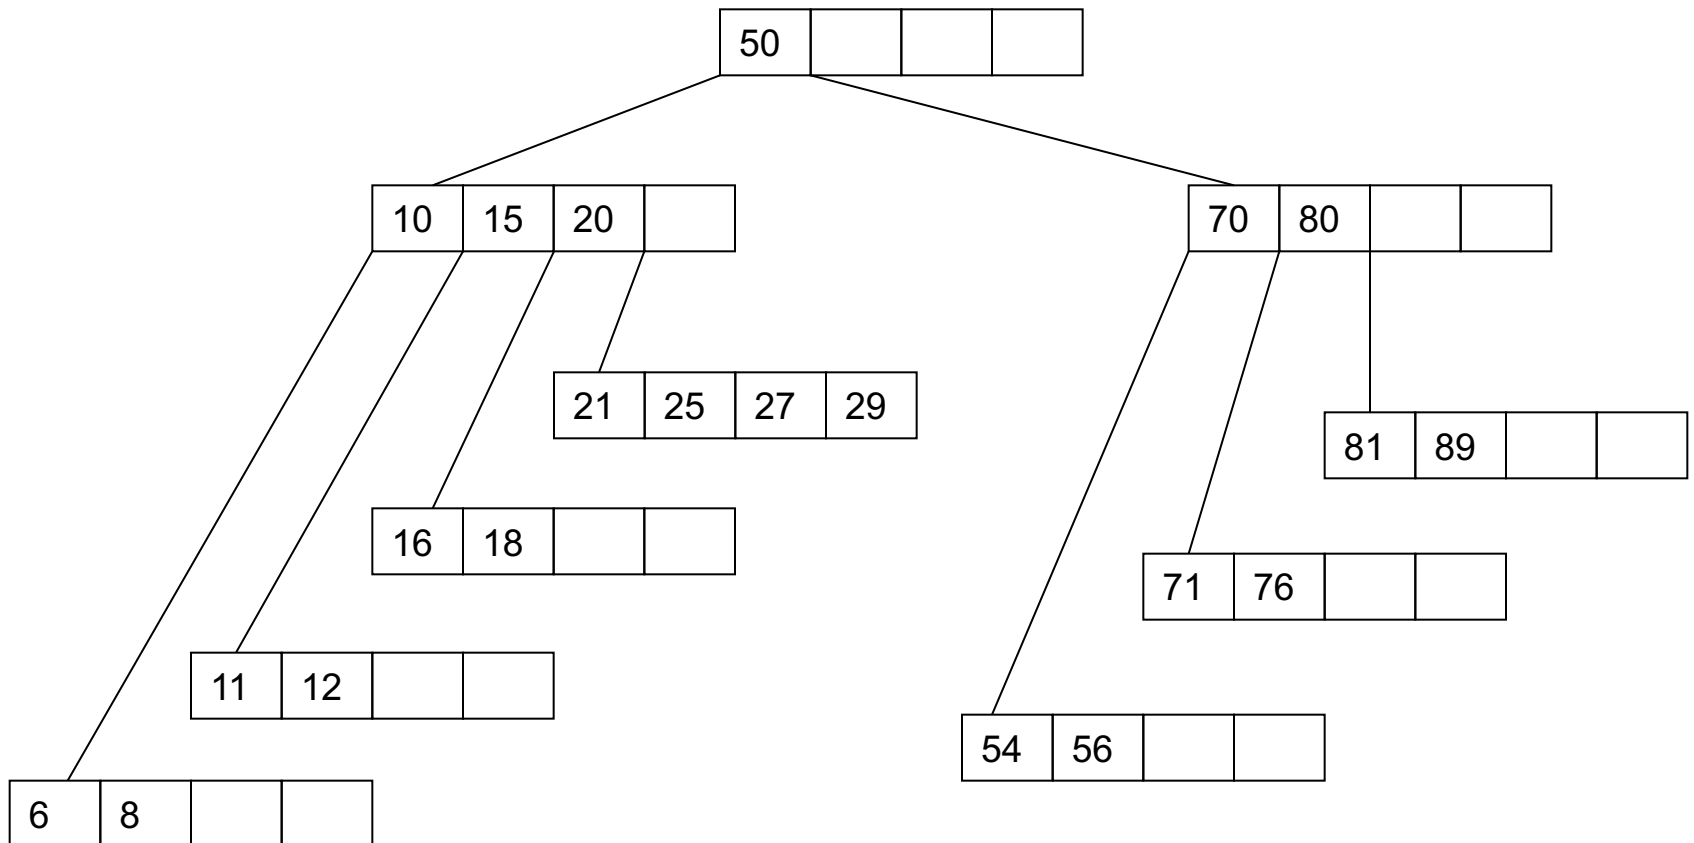
B-trees

```
1.  template <class T, int M>
2.  class BTreeNode {
3.  public:
4.      BTreeNode();
5.      BTreeNode(const T&);
6.  private:
7.      bool leaf;
8.      int keyTally;
9.      T keys[M-1];
10.     BTreeNode *pointers[M];
11.     friend BTree<T,M>;
12. }
```

Searching in a B-Tree



Searching



Searching

- ▶ An algorithm for finding a key in a B-tree is simple, and is coded as follows:

```
1. BTreeNode *BTreeSearch(keyType K, BTreeNode *node) {
2.     if (node != 0) {
3.         for (i = 1; i <= node->keyTally &&
4.              node->key[i-1] < K; i++);
5.         if (i > node->keyTally || node->keys[i-1] > K)
6.             return BTreeSearch(K, node->pointers[i-1]);
7.         else
8.             return node;
9.     }
```

Inserting a Key into a B-Tree



Inserting

- ▶ Both the **insertion and deletion** operations appear to be somewhat **challenging** if we remember that all **leaves** have to be **at the last level**.
- ▶ Implementing insertion becomes easier when the **strategy** of building a tree is **changed**.

Inserting

- ▶ When **inserting** a node into a binary search tree, the tree is always built from **top to bottom**, resulting in unbalanced trees.
- ▶ If the first incoming key is the **smallest**, then this key is put in the root, and the root does **not have a left subtree** unless special provisions are made to balance the tree.

Inserting

- ▶ But a tree can be built from the **bottom up** so that the root is an entity always in flux, and only at the end of all insertions can we know for sure the contents of the root.
- ▶ This strategy is applied to inserting keys into B-trees.

Inserting


- ▶ In the process,
 - ▶ Given an incoming key, we go **directly to a leaf** and place it there, if there is **room**.
 - ▶ When the **leaf is full**, another **leaf is created**, the keys are **divided between these leaves**, and one key is **promoted to the parent**.
 - ▶ If the **parent is full**, the process is **repeated** until the root is reached and a new root created.

Inserting

- ▶ To approach the problem more systematically, there are **three common situations** encountered when inserting a key into a B-tree.


Inserting

- I. A key is placed in a leaf that still **has some room**.




Inserting

2. The leaf in which a key should be placed is **full**.



Inserting

2. The leaf in which a key should be placed is **full**.



Inserting

2. The leaf in which a key should be placed is **full**.
 - In this case, the leaf is **split**, creating a **new leaf**, and **half** of the keys are **moved** from the full leaf to the new leaf.
 - But the new leaf has to be incorporated in the B-tree. The **middle key** is moved to the **parent**, and a **pointer** to the new leaf is placed in the parent as well.

Inserting

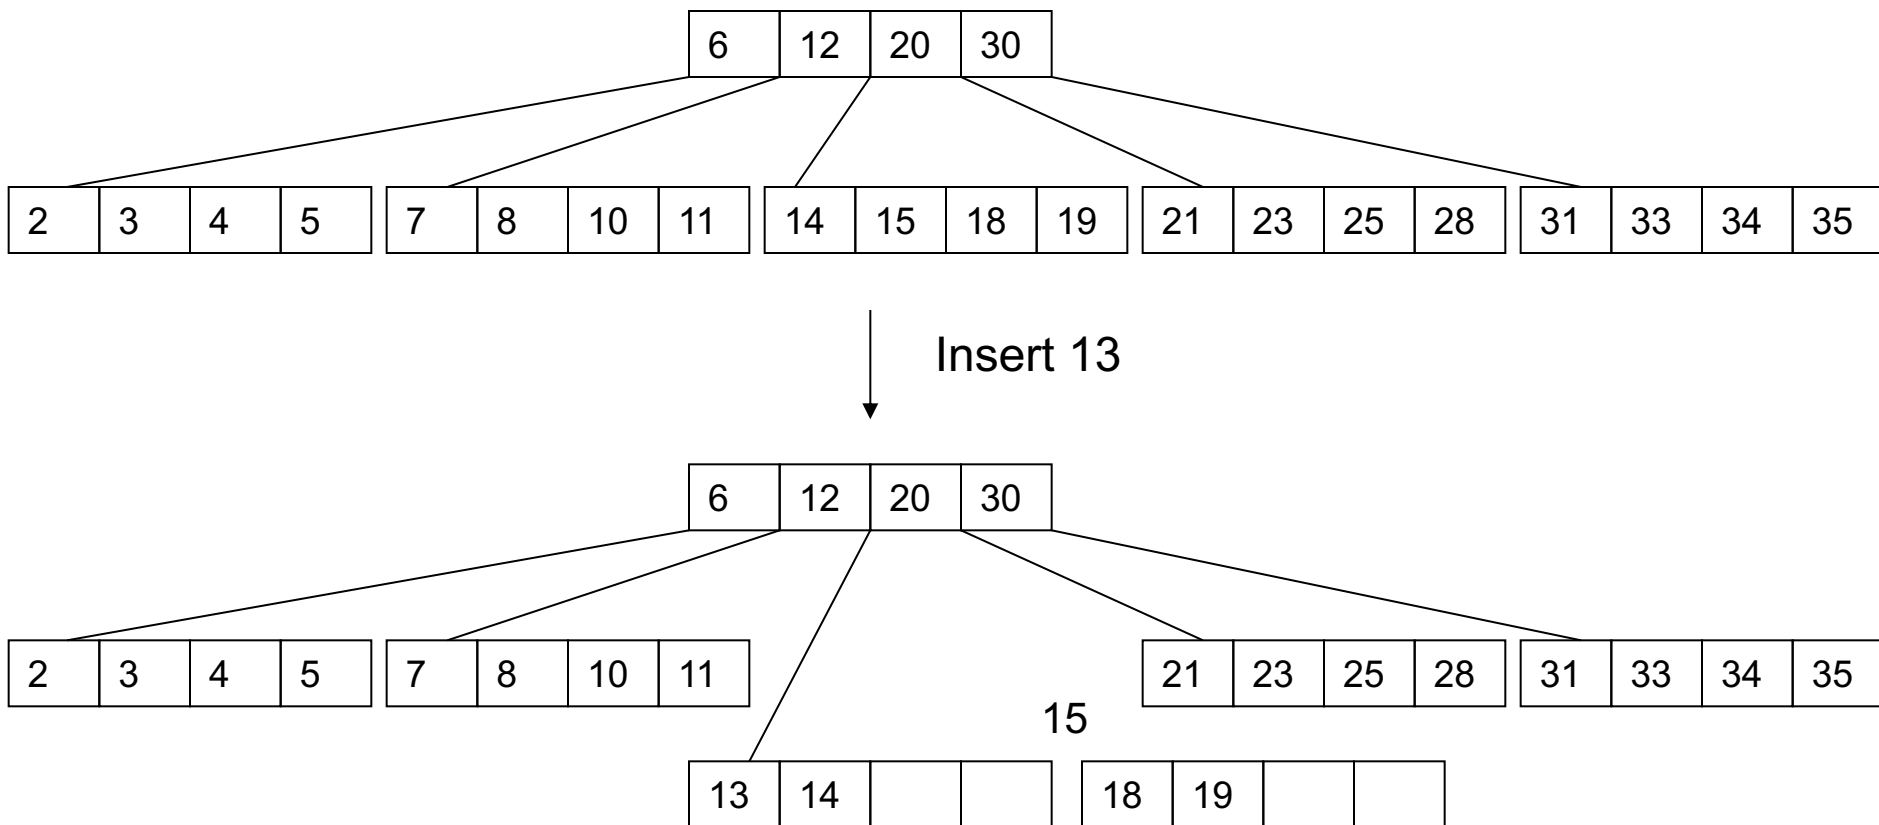
2. The leaf in which a key should be placed is **full**.
 - The same procedure can be **repeated** for each internal node of the B-tree so that each such split adds one more node to the B-tree.
 - Moreover, such a **split guarantees** that each **leaf** never has **less than $\lceil m/2 \rceil - 1$ keys**.

Inserting

3. A special case arises if the root of the B-tree is full.
 - In this case, a **new root** and a **new sibling** of the existing root have to be created.
 - This split results in **two new nodes** in the B-tree.



Inserting

3. A special case arises if the root of the B-tree is full.



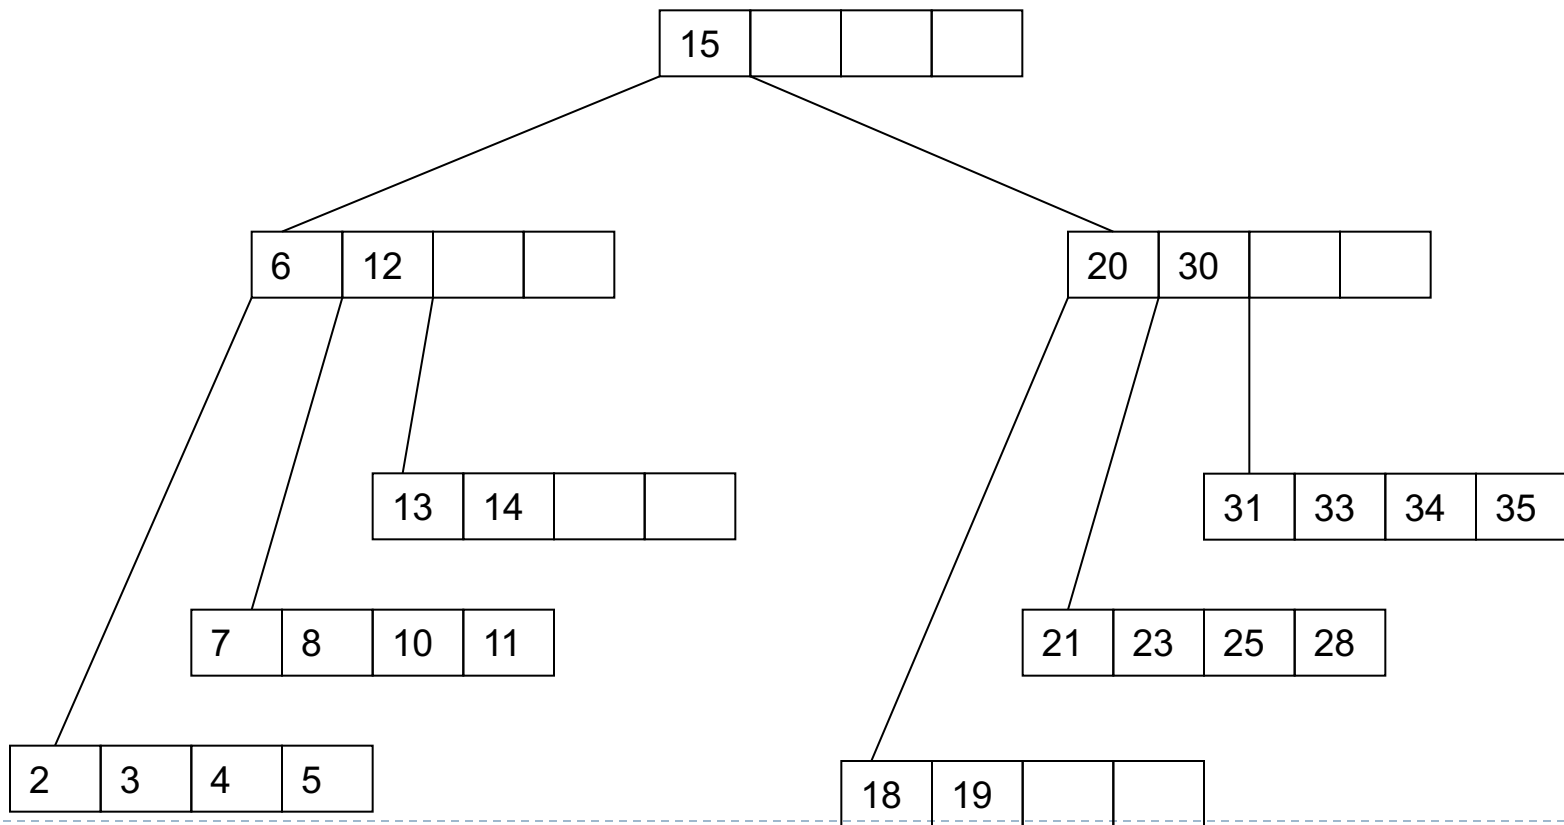
Inserting

3. A special case arises if the root of the B-tree is full.



Inserting

3. A special case arises if the root of the B-tree is full.



Deleting a Key from a B-Tree



Deleting

- ▶ Deletion is to a great extent a **reversal of insertion**, although it has **more special cases**.
- ▶ Care has to be taken to **avoid** allowing any node to be **less than half full** after a deletion.
- ▶ This means that nodes sometimes have to be **merged**.

Deleting

- ▶ There are two main cases:
 - ▶ Deleting a key from a **leaf** and
 - ▶ Deleting a key from a **nonleaf** node.
- ▶ In the latter case, we will use a procedure similar to **delete by copying** used for binary search trees.

Deleting

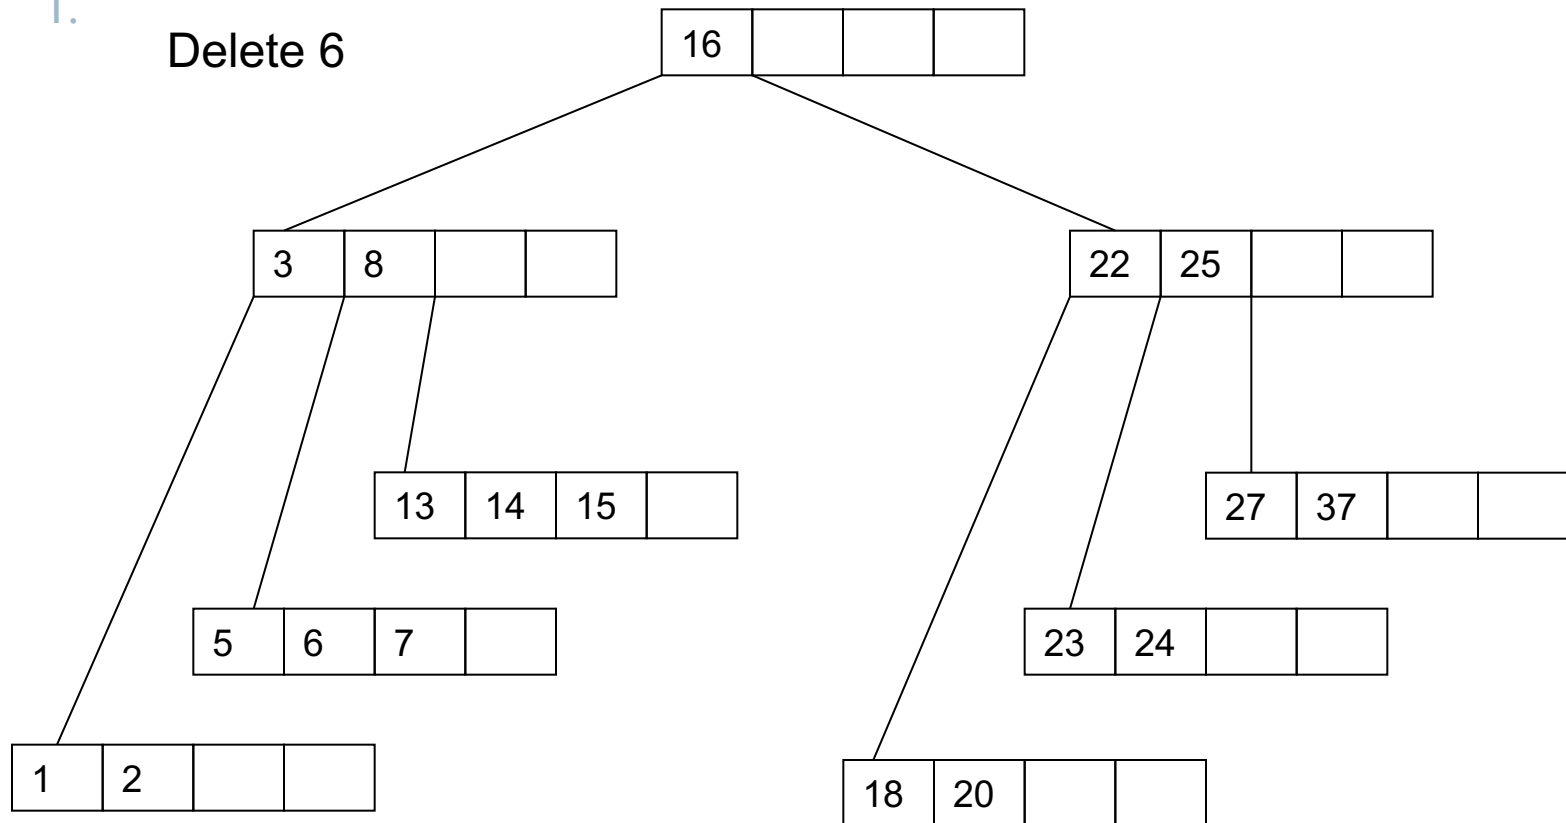
- ▶ Deleting a key from a **leaf**
 - I. If, after deleting a key K , the leaf is **at least half full** and only keys greater than K are moved to the left to fill the hole
 - This is the inverse of insertion's case I.

Deleting

► Deleting a key from a **leaf**

I.

Delete 6

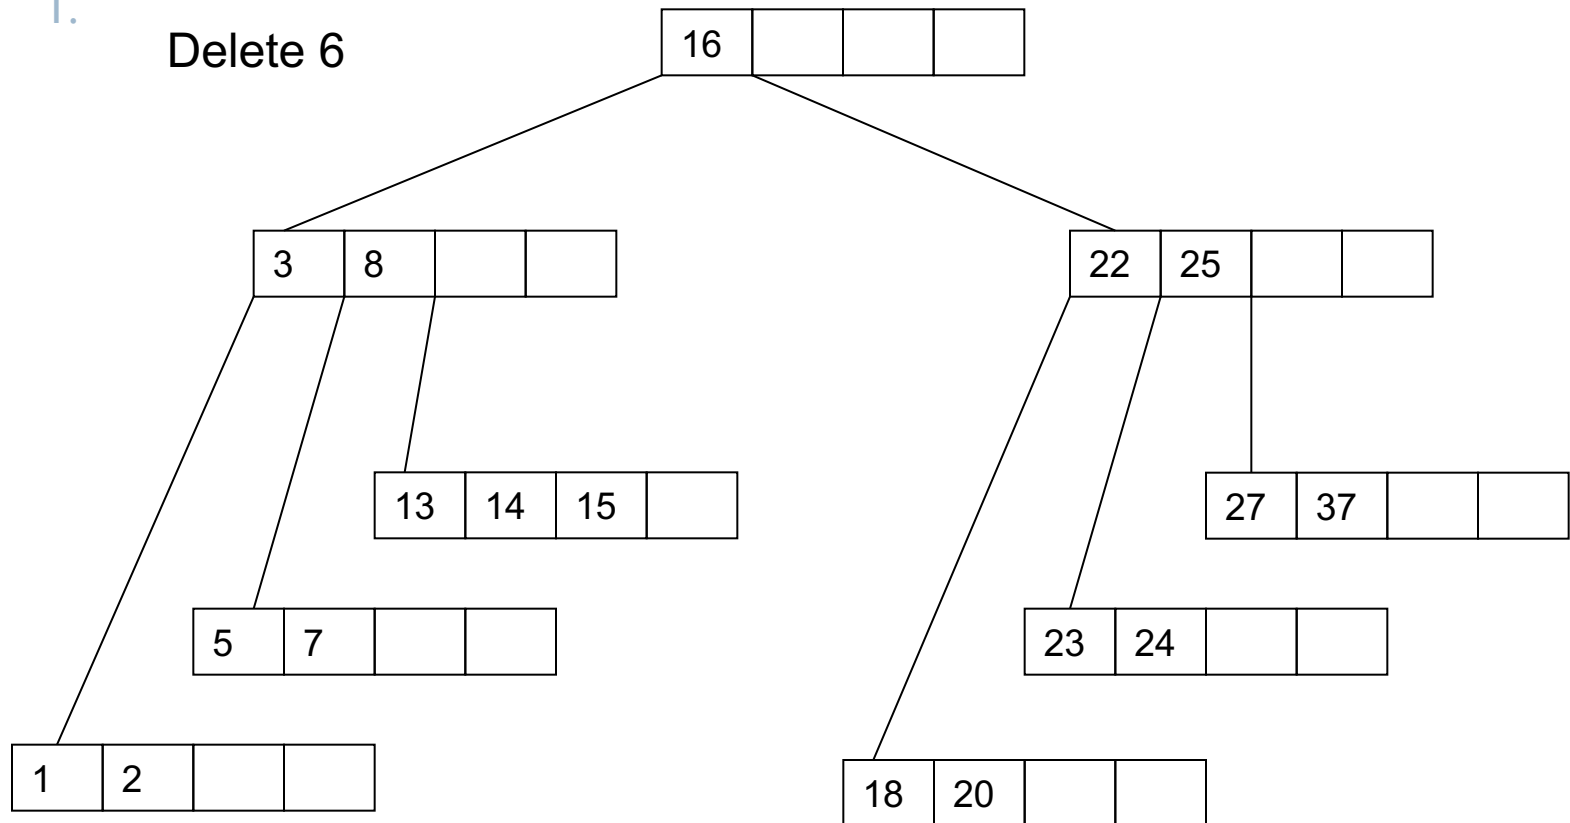


Deleting

► Deleting a key from a **leaf**

I.

Delete 6



Deleting

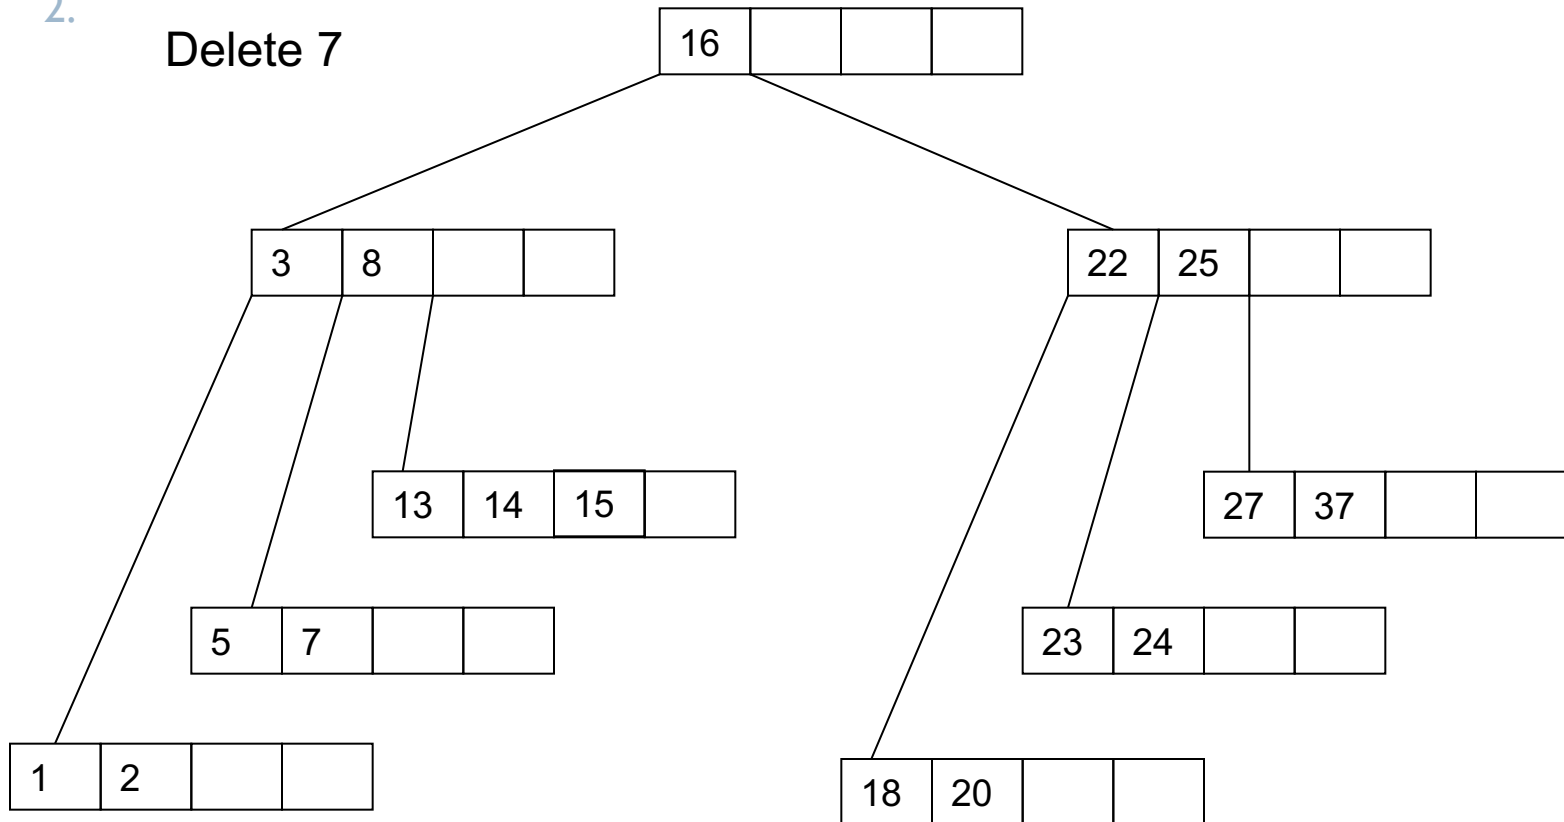
- ▶ Deleting a key from a **leaf**
- 2. If, after deleting K , the number of keys in the leaf is **less than $\lceil m/2 \rceil - 1$** , causing an **underflow**:
 - 1. If there is a **left or right sibling** with the number of keys **exceeding the minimal $\lceil m/2 \rceil - 1$** , then
 - **all keys** from this leaf and this sibling are **redistributed** between them by
 - moving the **separator key** from the **parent to the leaf** and
 - moving the **middle key** from the node and the sibling combined **to the parent**.

Deleting

- ▶ Deleting a key from a **leaf**

2.

Delete 7

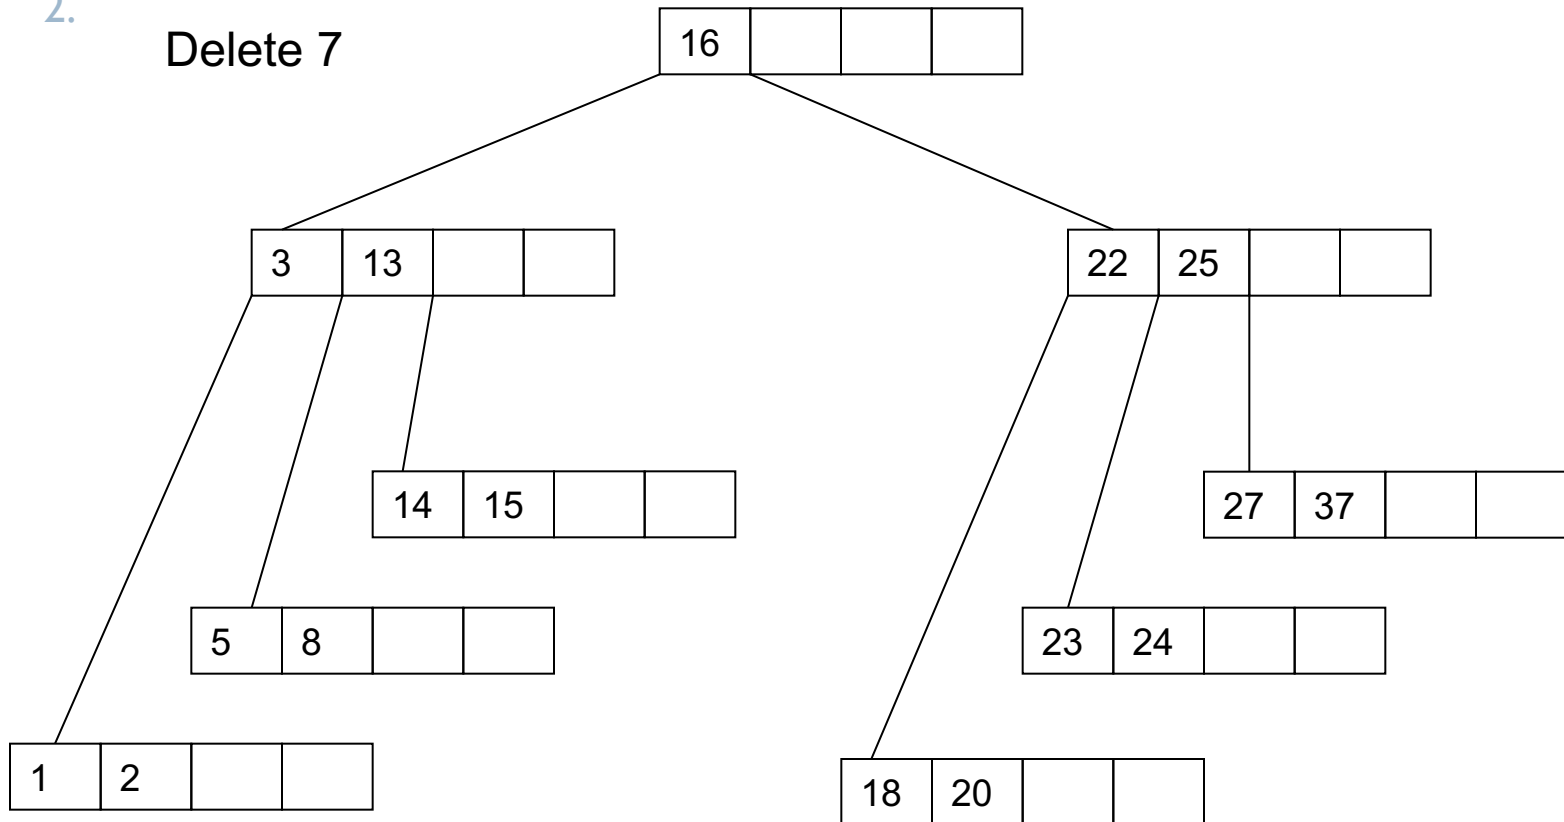


Deleting

- ▶ Deleting a key from a **leaf**

2.

Delete 7



Deleting

- ▶ Deleting a key from a **leaf**
- 2. If, after deleting K , the number of keys in the leaf is **less than $\lceil m/2 \rceil - 1$** , causing an **underflow**:
- 2. If the leaf underflows and the number of keys in its siblings is $\lceil m/2 \rceil - 1$,
 - then the leaf and a sibling are **merged**;
 - the key from the **leaf**, from its **sibling**, and the **separating key** from the parent are all **put in the leaf**, and
 - the sibling node is **discarded**.

Deleting

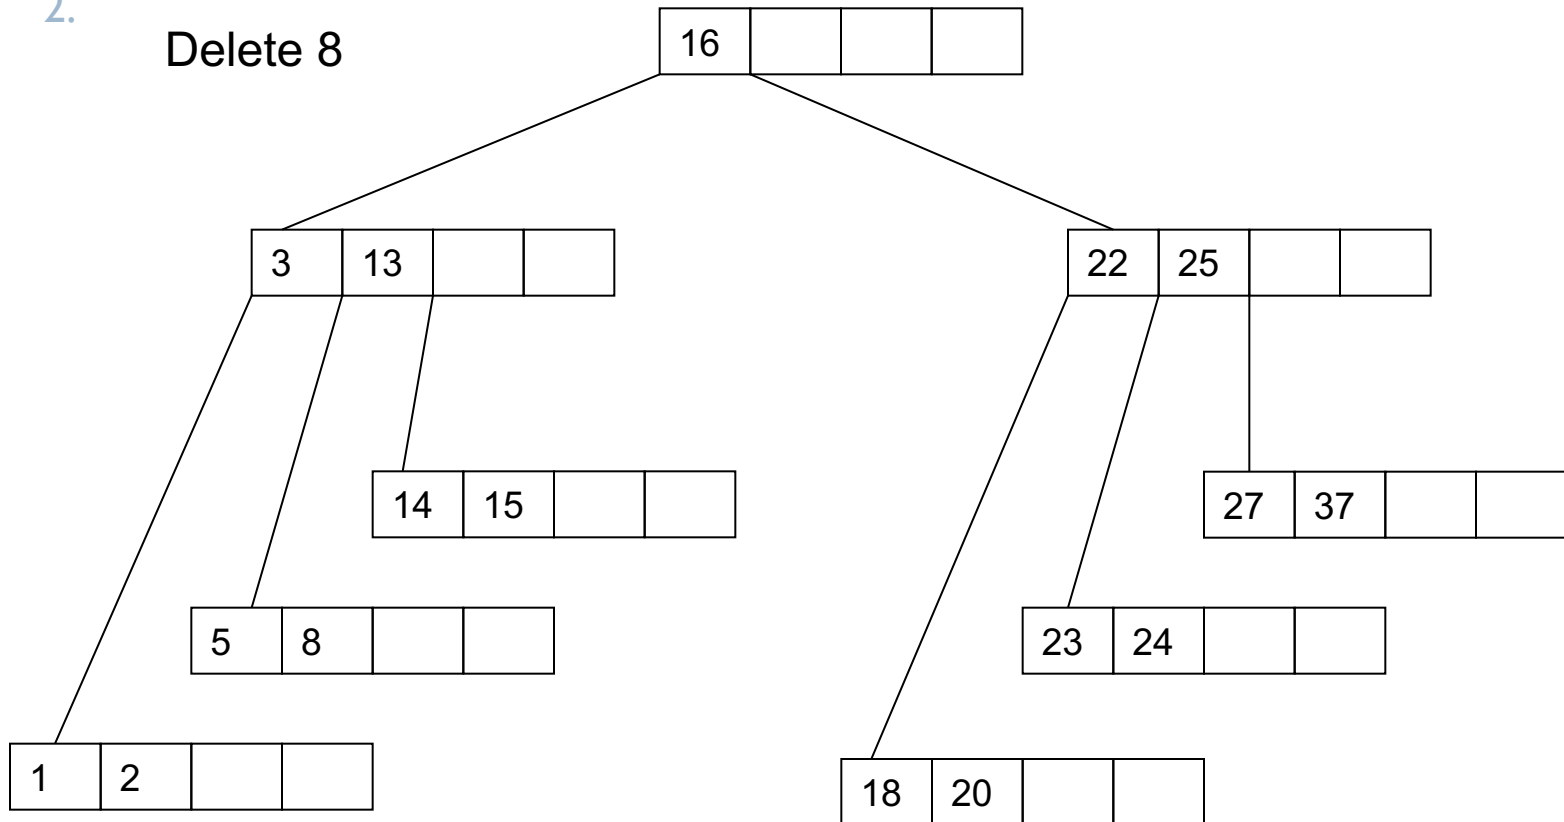
- ▶ Deleting a key from a **leaf**
- 2. If, after deleting K , the number of keys in the leaf is **less than $\lceil m/2 \rceil - 1$** , causing an **underflow**:
 - 2. If the leaf underflows and the number of keys in its siblings is $\lceil m/2 \rceil - 1$,
 - The keys in the parent are **moved** if a **hole** appears.
 - This can initiate a **chain of operations** if the parent **underflows**. The parent is now treated as though it were a leaf, and either step 2.2 is **repeated** until step 2.1 can be executed or the root of the tree has been reached.
 - This is the inverse of insertion's case 2.

Deleting

- ▶ Deleting a key from a **leaf**

2.

Delete 8

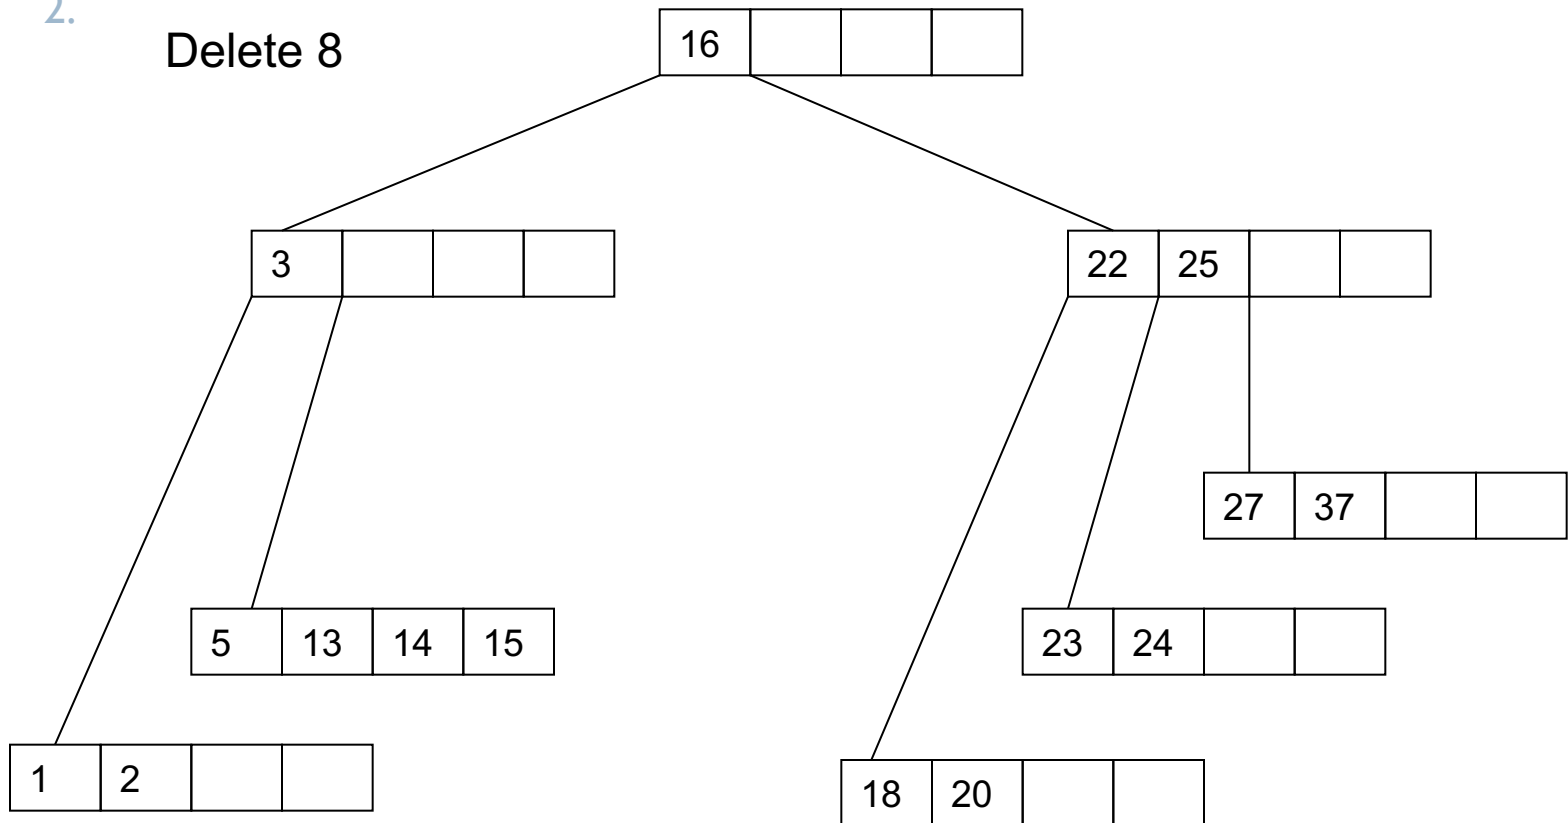


Deleting

- ▶ Deleting a key from a **leaf**

2.

Delete 8



Deleting

- ▶ Deleting a key from a **leaf**
- 2. If, after deleting K , the number of keys in the leaf is **less than $\lceil m/2 \rceil - 1$** , causing an **underflow**:
 - 2. A **particular case** results in merging a leaf or nonleaf with its sibling when its **parent is the root with only one key**.
 - In this case, the keys from the **node** and its **sibling**, along with the only key of the **root**, are **put in the node**, which becomes a **new root**, and
 - both the **sibling** and the **old root** nodes are **discarded**.

Deleting

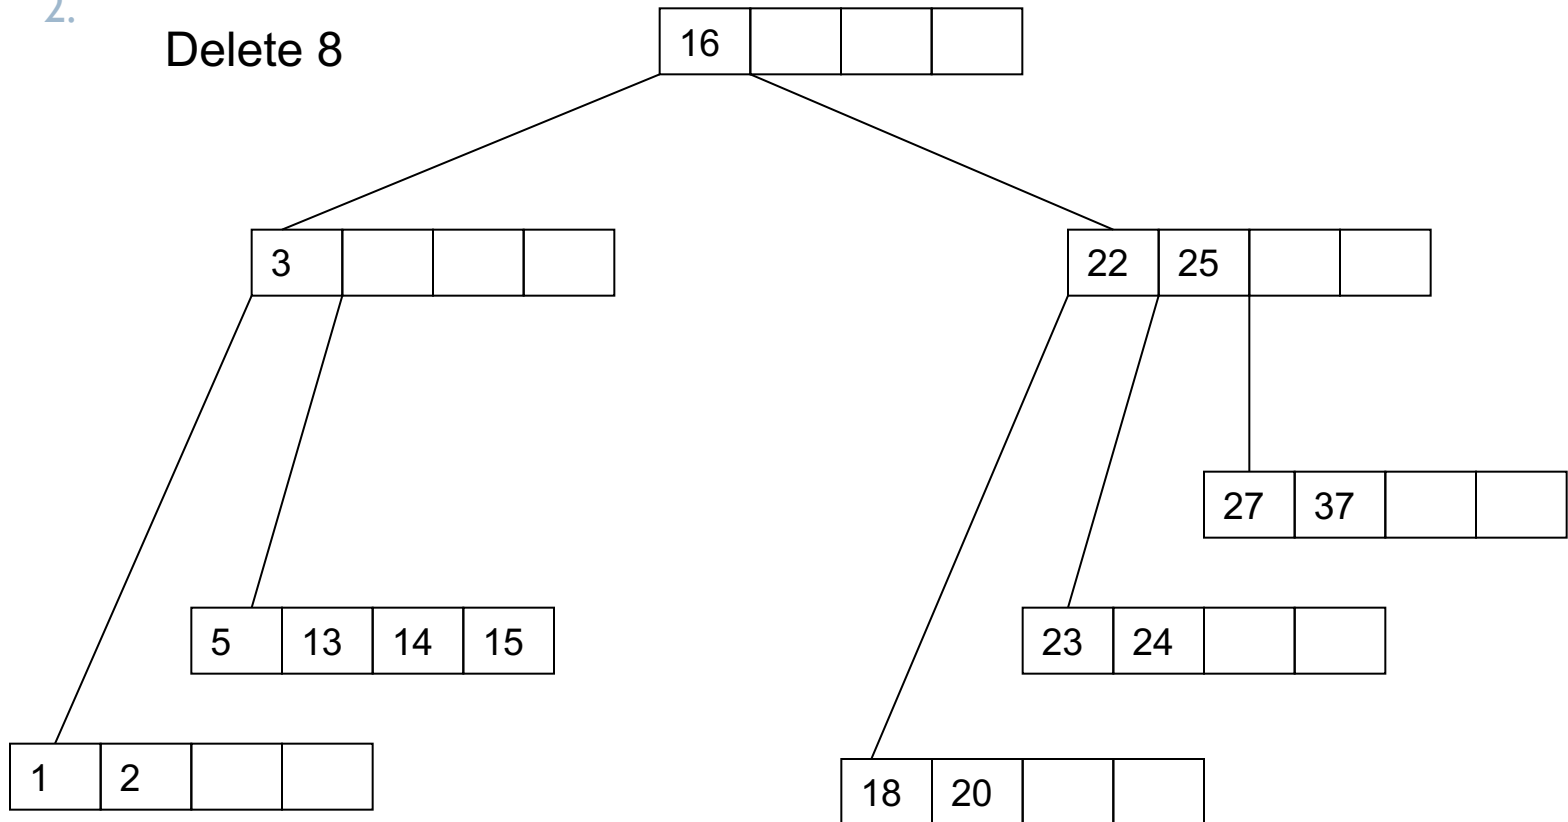
- ▶ Deleting a key from a **leaf**
- 2. If, after deleting K , the number of keys in the leaf is **less than $\lceil m/2 \rceil - 1$** , causing an **underflow**:
 - 2. A **particular case** results in merging a leaf or nonleaf with its sibling when its **parent is the root with only one key**.
 - This is the only case when **two nodes disappear** at one time.
 - Also the **height** of the tree is **decreased by one**.
 - This is the inverse of insertion's case 3.

Deleting

- ▶ Deleting a key from a **leaf**

2.

Delete 8

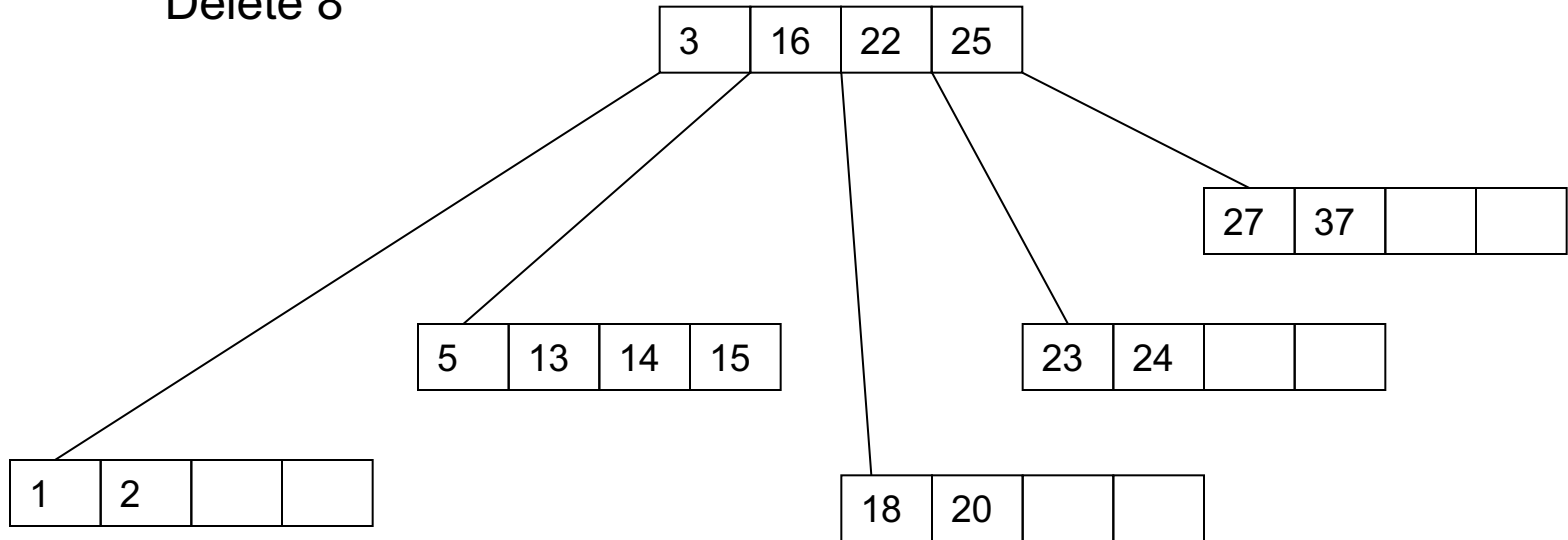


Deleting

- ▶ Deleting a key from a **leaf**

2.

Delete 8



Deleting

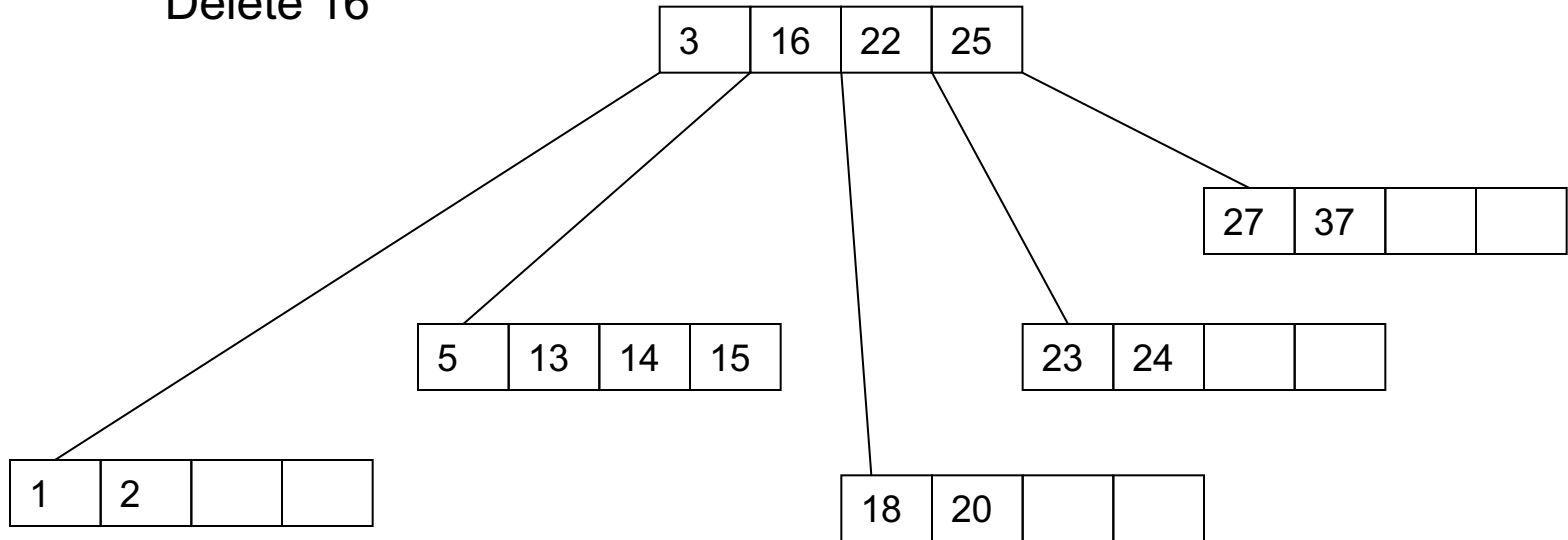
- ▶ Deleting a key from a **nonleaf**
 - This may lead to problems with **tree reorganization**.
 - Therefore, **deletion** from a **nonleaf node** is reduced to deleting a key from a **leaf**.
 - The key to be deleted is **replaced by its immediate predecessor** (the successor could also be used), which can only be found in a leaf.
 - This **successor key is deleted** from the leaf, which brings us to the preceding case I

Deleting

- ▶ Deleting a key from a **nonleaf**

2.

Delete 16




Look for its min. key of the right sub-tree -> 18

Deleting

► Deleting a key from a **nonleaf**

2.

Delete 16




Replace 16 by 18 and remove 18 at the leaf node

Deleting

► Deleting a key from a **nonleaf**

2.

Delete 16




Since the leaf node is in short and it can't borrow from its right sibling, it will borrow a key from its left sibling. Pull 18 down and push 15 up.

B+ Tree and its Properties

- ▶ A B+ tree of **order $M \geq 3$** is an M-ary tree with the following properties:
 1. The **root** is either a **leaf** or has between **1** and **$M-1$** keys
 2. Each **node**, except the root, has between **$\text{ceil}(M/2) - 1$** and **$M-1$** keys
 3. Each **node** has between **$\text{ceil}(M/2)$** and **M** children
 4. The keys at each node are **ordered**
 5. The **data items** are stored at the **leaves**
 - ▶ All leaves are at the same depth
 - ▶ Each leaf has between $\text{ceil}(L/2)$ and L data items, for some L (L is much smaller than M in general, but we will assume $M = L$ in most examples)

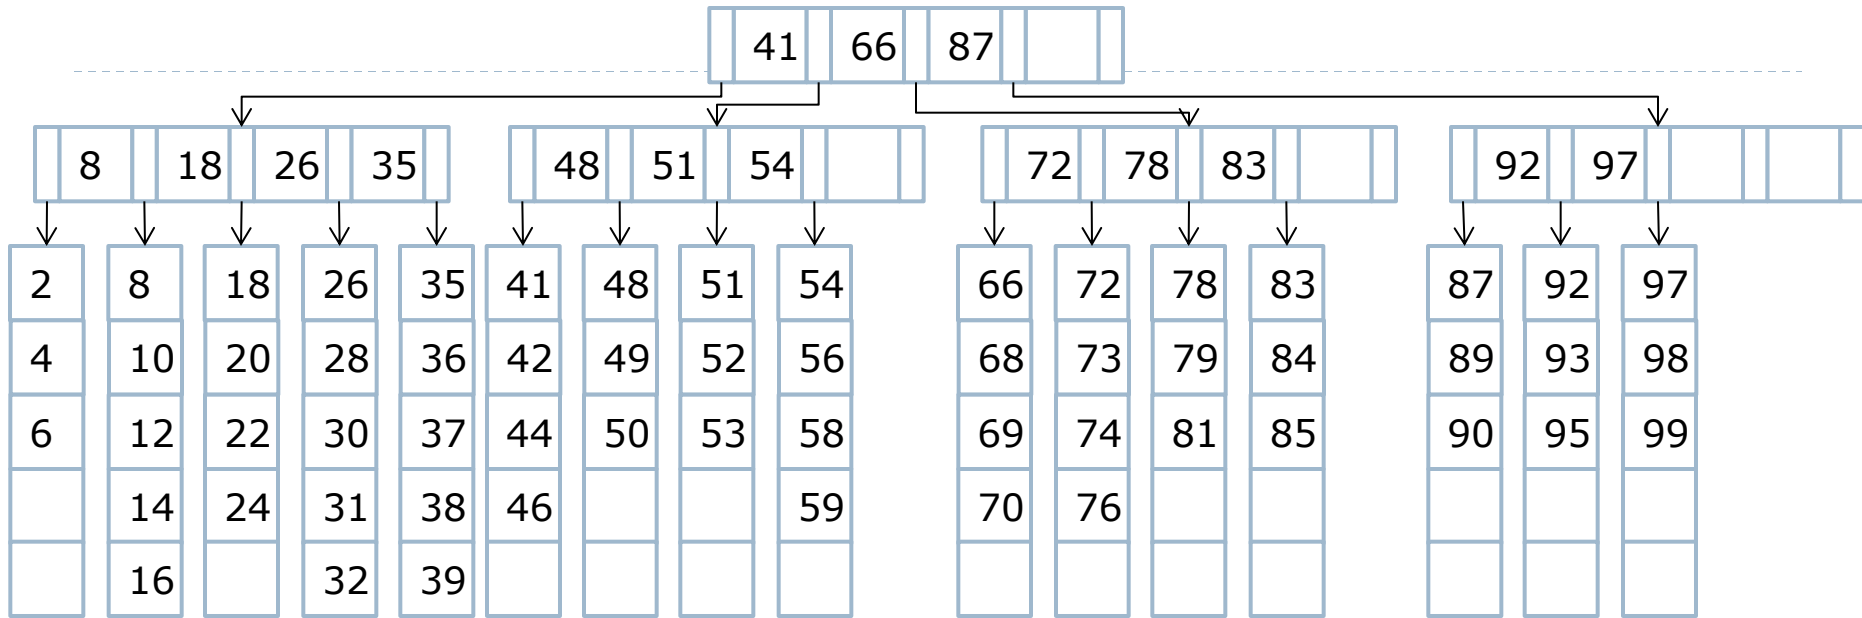
B+ Tree – Example



A B+ Tree of order 3

- ▶ The root is either a leaf or has between 1 and $3 - 1 = 2$ keys
- ▶ Each node, except the root, has between $\text{ceil}(3/2) - 1 = 1$ and $3 - 1 = 2$ keys
- ▶ Each node has at between $\text{ceil}(3/2) = 2$ and 3 children
- ▶ The keys at each node are ordered
- ▶ The data items are stored at the leaves
 - ▶ All leaves are at the same depth
 - ▶ Each leaf has between $\text{ceil}(3/2) = 2$ and 3 data items

B+ Tree – Example



A B+ Tree of order 5

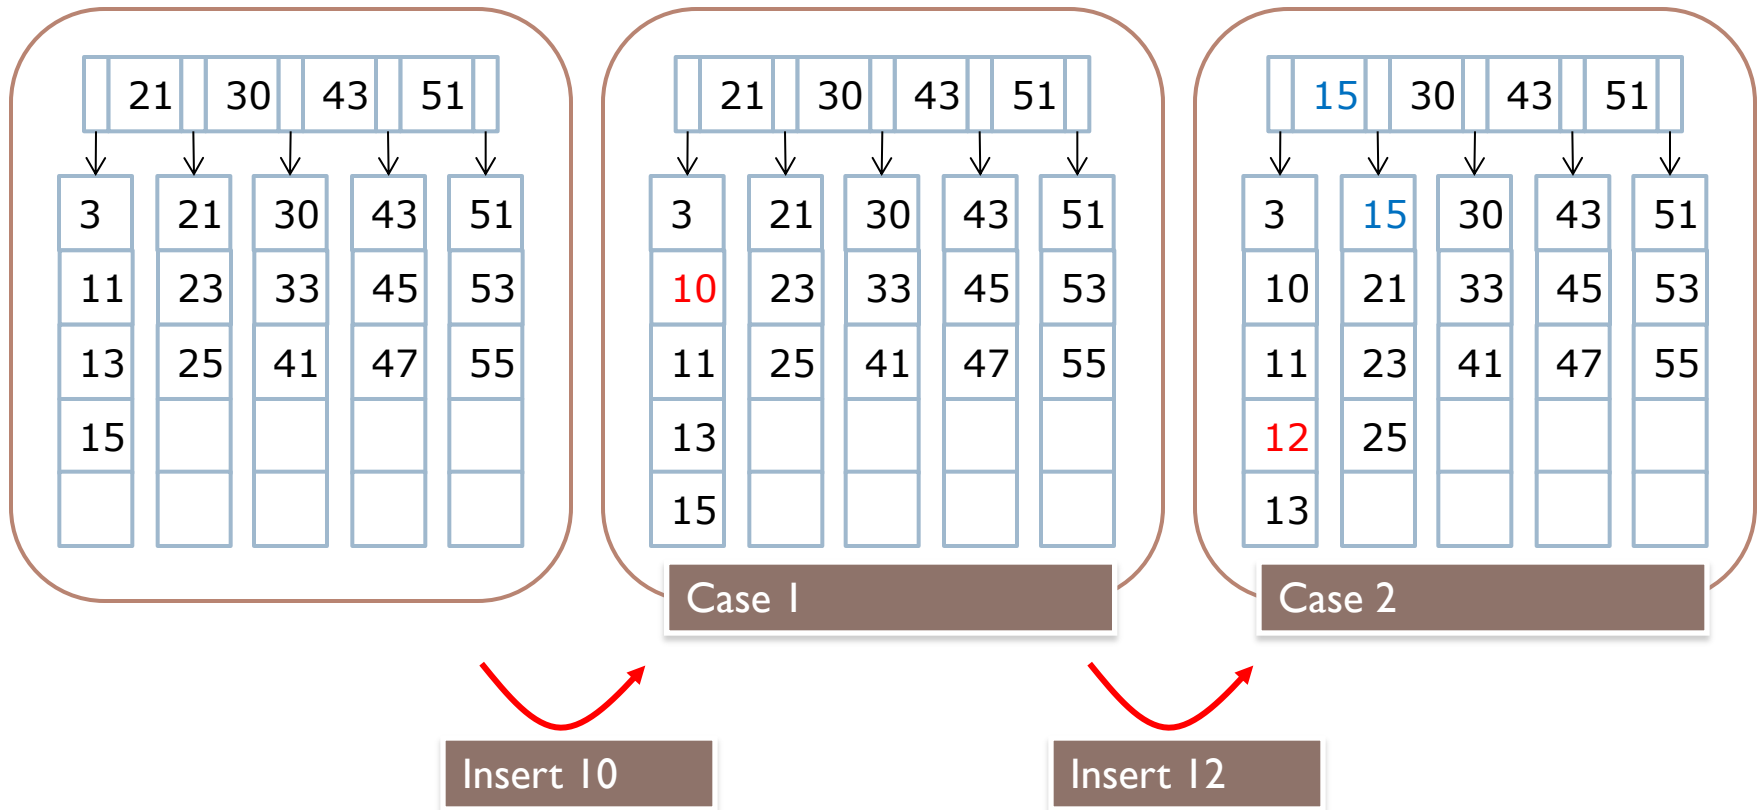
- ▶ The root is either a leaf or has between 1 and $5 - 1 = 4$ keys
- ▶ Each node, except the root, has between $\text{ceil}(5/2) - 1 = 2$ and $5 - 1 = 4$ keys
- ▶ Each node has at between $\text{ceil}(5/2) = 3$ and 5 children
- ▶ The keys at each node are ordered
- ▶ The data items are stored at the leaves
 - ▶ All leaves are at the same depth
 - ▶ Each leaf has between $\text{ceil}(5/2) = 3$ and 5 data items

B+ Tree – Insertion


- ▶ To insert an element to a B+ tree, **search** the tree using a search algorithm similar to the one for Binary Search Tree and find the leaf node, say n , where the new item should be inserted
 1. If the leaf has enough space, insert the new item in the node n in sorted order
 2. Else if sibling node, say p , has space
 - a) Insert the new item in the node n in sorted order
 - b) Move the last item of n into p
 - c) Update the corresponding key in the parent node so that first key of p is in the parent
 3. Otherwise,
 - a) Split n into n and a new node n_2
 - b) Re-distribute elements evenly into n and n_2
 - c) Insert key of the first element of n_2 into parent of n
 - d) If parent of n has keys less than $M-1$, done!
 - e) Else repeat a) to d)

B+ Tree – Insertion Example


A B+ Tree of order 5



B+ Tree – Insertion Example




Insert 27



Case 1

Insert 9



Case 3

B+ Tree – Deletion

- ▶ To delete an element, say k , in a B+ tree, search the tree using a search algorithm similar to the one for Binary Search Tree and find the leaf node, say n , where the item should be removed
 - I. If n is root, remove k
 - ▶ If root has more than one keys, done!!!
 - ▶ If root has only k
 - If any of its child node can lend a node, borrow key from the child and adjust child links
 - ▶ Otherwise, merge the children nodes (it will be a new root)

B+ Tree – Deletion


2. If n is other node, remove k

- ▶ If n has at least $\text{ceil}(M/2)$ keys, done!!!
- ▶ If n has less than $\text{ceil}(M/2) - 1$ keys,
 - If a sibling can lend a key, borrow key from the sibling and adjust keys in n and the parent node, adjust child links
 - Else, merge n with its sibling and adjust child links

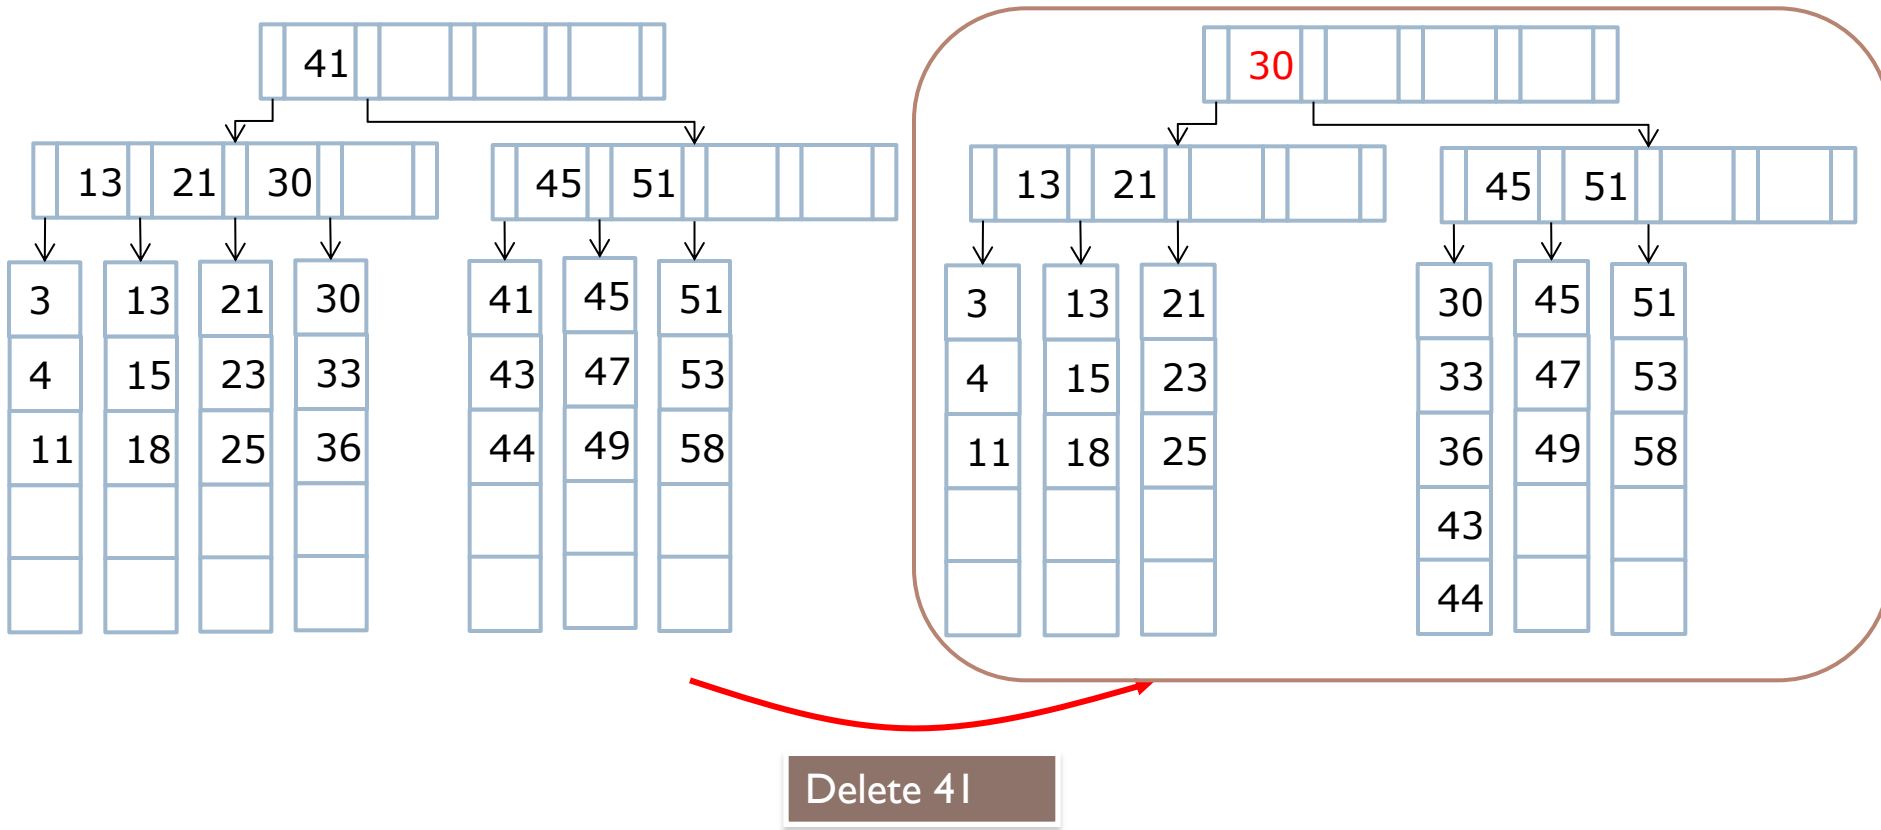
3. If n is a leaf node, remove k

- ▶ If n has at least $\text{ceil}(M/2)$ elements, done!!! (In case the smallest key is deleted, push up the next key)
- ▶ If n has less than $\text{ceil}(M/2)$ elements
 - If the sibling can lend a key, borrow key from a sibling and adjust keys in n and the parent node
- ▶ Else, merge n and its sibling and adjust keys in the parent node



B+ Tree – Deletion Example



B+ Tree – Deletion Example



B+ Tree – Deletion Example



Delete 3

CHAPTER 7 END