

--	--	--	--	--	--	--	--

Day: ☐ Monday ☐ TuesdayTime: ☐ 10:00 - 10:50 ☐ 11:00 - 11:50 ☐ 12:00 - 12:50 ☐ 14:00 - 14:50 ☐ 18:00 - 18:50

The Process

Introduction

Topics to be covered in this tutorial include:

1. How process states change as programs run and either use the CPU (e.g., perform an add instruction) or do I/O (e.g., send a request to a disk and wait for it to complete).

Points to note about the simulator:

1. This tutorial is based on running a **simulator**, which mimics some aspect of an operating system.
2. The simulator can be used to both **generate problems** and **obtain solutions** for an infinite number of problems. Different random seeds can usually be used to generate different problems; using the `-c` flag computes the answers for you (presumably after you have tried to compute them yourself!).

Acknowledgement

This tutorial was adapted from OSTEP book written by Remzi and Andrea Arpaci-Dusseau at University of Wisconsin. This free OS book is available at <http://www.ostep.org>.

Pre-lab

1. Logging in to the Linux server

- Start the SSH client, e.g., MobaXterm or Xshell.
- Log in to the Linux server using the following details:

Host Name: gateway.cs.cityu.edu.hk

User Name: your EID (e.g., cctom2)

Password: your password

⚠ Your password will not be shown on the screen as you type it, not even as a row of stars (*****).

NOTE: The shell will always give you a prompt if it is ready to accept commands. The shell prompt normally ends in a \$ sign as we use in this tutorial. Some shell prompts end in % or > instead. Never copy/type the shell prompt used in this tutorial.

NOTE: Please don't forget to log out (use the `exit` command) after you finish your work.

2. Getting the `process-run.py` simulator

A simulation program, called `process-run.py`, allows you to see how the state of a process state changes as it runs on a CPU. This simulator is located in the directory `/public/cs3103/tutorial2`.

Start by copying the simulator to a directory in which you plan to do your work. For example, to copy the file named `process-run.py` in the directory `/public/cs3103/tutorial2` to your current directory, enter:

```
$ cp /public/cs3103/tutorial2/process-run.py .
```

The last dot/period (.) indicates the current directory as destination.

Introduction to process-run.py

This program, `process-run.py`, allows you to see how process states change as programs run and either use the CPU (e.g., perform an add instruction) or do I/O (e.g., send a request to a disk and wait for it to complete).

As described in the lecture, processes can be in a few different states:

RUNNING: the process is using the CPU right now

READY: the process could be using the CPU right now but some other process is

WAITING: the process is waiting on I/O (e.g., it issued a request to a disk)

DONE: the process is finished executing

In this tutorial, we'll see how these process states change as a program runs, and thus learn a little bit better how these things work.

To run the program and get its options, do this:

```
$ ./process-run.py -h
```

If this doesn't work, type "python2" before the command, like this:

```
$ python2 process-run.py -h
```

What you should see is this:

```
Usage: process-run.py [options]
```

Options:

```
-h, --help          show this help message and exit
-s SEED, --seed=SEED the random seed
-l PROCESS_LIST, --processlist=PROCESS_LIST
                    a comma-separated list of processes to run, in the
                    form X1:Y1,X2:Y2,... where X is the number of
                    instructions that process should run, and Y the
                    chances (from 0 to 100) that an instruction will use
                    the CPU or issue an IO
-L IO_LENGTH, --iolength=IO_LENGTH
                    how long an IO takes
-S PROCESS_SWITCH_BEHAVIOR, --switch=PROCESS_SWITCH_BEHAVIOR
                    when to switch between processes: SWITCH_ON_IO,
                    SWITCH_ON_END
```

```
-I IO_DONE_BEHAVIOR, --iodone=IO_DONE_BEHAVIOR
                        type of behavior when IO ends: IO_RUN_LATER,
                        IO_RUN_IMMEDIATE
-c                      compute answers for me
-p, --printstats       print statistics at end; only useful with -c flag
                        (otherwise stats are not printed)
```

The most important option to understand is the `PROCESS_LIST` (as specified by the `-l` or `--processlist` flags) which specifies exactly what each running program (or "process") will do. A process consists of instructions, and each instruction can just do one of two things:

- use the CPU
- issue an IO (and wait for it to complete)

When a process uses the CPU (and does no IO at all), it should simply alternate between **RUNNING** on the CPU or being **READY** to run. For example, here is a simple run that just has one program being run, and that program only uses the CPU (it does no IO).

```
$ ./process-run.py -l 5:100
```

Produce a trace of what would happen when you run these processes:

Process 0

```
cpu
cpu
cpu
cpu
cpu
```

Important behaviors:

System will switch when the current process is `FINISHED` or `ISSUES AN IO`
After IOs, the process issuing the IO will run `LATER` (when it is its turn)

Here, the process we specified is "5:100" which means it should consist of 5 instructions, and the chances that each instruction is a CPU instruction are 100%.

You can see what happens to the process by using the `-c` flag, which computes the answers for you:

```
$ ./process-run.py -l 5:100 -c
```

Time	PID: 0	CPU	IOs
1	RUN:cpu	1	
2	RUN:cpu	1	
3	RUN:cpu	1	
4	RUN:cpu	1	
5	RUN:cpu	1	

This result is not too interesting: the process is simple in the **RUN** state and then finishes, using the CPU the whole time and thus keeping the CPU busy the entire run, and not doing any I/Os.

Let's make it slightly more complex by running two processes:

```
$ ./process-run.py -l 5:100,5:100
```

Produce a trace of what would happen when you run these processes:

Process 0

cpu

cpu

cpu

cpu

cpu

Process 1

cpu

cpu

cpu

cpu

cpu

Important behaviors:

Scheduler will switch when the current process is **FINISHED** or **ISSUES AN IO**

After I/Os, the process issuing the IO will run **LATER** (when it is its turn)

In this case, two different processes run, each again just using the CPU. What happens when the operating system runs them? Let's find out:

```
$ ./process-run.py -l 5:100,5:100 -c
```

Time	PID: 0	PID: 1	CPU	I/Os
1	RUN:cpu	READY	1	
2	RUN:cpu	READY	1	
3	RUN:cpu	READY	1	
4	RUN:cpu	READY	1	
5	RUN:cpu	READY	1	
6	DONE	RUN:cpu	1	
7	DONE	RUN:cpu	1	
8	DONE	RUN:cpu	1	
9	DONE	RUN:cpu	1	
10	DONE	RUN:cpu	1	

As you can see above, first the process with "process ID" (or "PID") 0 runs, while process 1 is **READY** to run but just waits until 0 is done. When 0 is finished, it moves to the **DONE** state, while 1 runs. When 1 finishes, the trace is done.

Let's look at one more example before getting to some questions. In this example, the process just issues I/O requests. We specify here that I/Os take 5 time units to complete with the flag -L.

```
$ ./process-run.py -l 3:0 -L 5
```

Produce a trace of what would happen when you run these processes:

Process 0

io-start

io-start

io-start

Important behaviors:

System will switch when the current process is FINISHED or ISSUES AN IO

After IOs, the process issuing the IO will run LATER (when it is its turn)

What do you think the execution trace will look like? Let's find out:

```
$ ./process-run.py -l 3:0 -L 5 -c
```

Time	PID: 0	CPU	IOs
1	RUN:io-start	1	
2	WAITING		1
3	WAITING		1
4	WAITING		1
5	WAITING		1
6*	RUN:io-start	1	
7	WAITING		1
8	WAITING		1
9	WAITING		1
10	WAITING		1
11*	RUN:io-start	1	
12	WAITING		1
13	WAITING		1
14	WAITING		1
15	WAITING		1
16*	DONE		

As you can see, the program just issues three I/Os. When each I/O is issued, the process moves to a **WAITING** state, and while the device is busy servicing the I/O, the CPU is idle.

Let's print some stats (run the same command as above, but with the -p flag) to see some overall behaviors:

```
Stats: Total Time 16
```

```
Stats: CPU Busy 3 (18.75%)
```

```
Stats: IO Busy 12 (75.00%)
```

As you can see, the trace took 16 clock ticks to run, but the CPU was only busy less than 20% of the time. The IO device, on the other hand, was quite busy. In general, we'd like to keep all the devices busy, as that is a better use of resources.

There are a few other important flags:

- **-s SEED, --seed=SEED**: the random seed. This gives you way to create a bunch of different jobs randomly.
- **-L IO_LENGTH, --iolength=IO_LENGTH**: this determines how long IOs take to complete (default is 5 ticks).
- **-S PROCESS_SWITCH_BEHAVIOR, --switch=PROCESS_SWITCH_BEHAVIOR**: when to switch between processes: SWITCH_ON_IO, SWITCH_ON_END. This determines when we switch to another process: SWITCH_ON_IO, the system will switch when a process issues an IO; SWITCH_ON_END, the system will only switch when the current process is done.
- **-I IO_DONE_BEHAVIOR, --iodone=IO_DONE_BEHAVIOR**: type of behavior when IO ends: IO_RUN_LATER, IO_RUN_IMMEDIATE. This determines when a process runs after it issues an IO: IO_RUN_IMMEDIATE: switch to this process right now; IO_RUN_LATER: switch to this process when it is natural to (e.g., depending on process-switching behavior).

Now go answer the questions to learn more.

Questions

1. Run `process-run.py` with the following flags: `-l 5:100,5:100`. What should the CPU utilization be (e.g., the percent of time the CPU is in use?) Why do you know this? Use the `-c` and `-p` flags to see if you were right.
2. Now run with these flags: `-l 4:100,1:0`. These flags specify one process with 4 instructions (all to use the CPU), and one that simply issues an I/O and waits for it to be done. How long does it take to complete both processes? Use `-c` and `-p` to find out if you were right.
3. Switch the order of the processes: `-l 1:0,4:100`. What happens now? Does switching the order matter? Why? (As always, use `-c` and `-p` to see if you were right)
4. We'll now explore some of the other flags. One important flag is `-S`, which determines how the system reacts when a process issues an I/O. With the flag set to `SWITCH_ON_END`, the system will **NOT** switch to another process while one is doing I/O, instead waiting until the process is completely finished. What happens when you run the following two processes (`-l 1:0,4:100 -c -S SWITCH_ON_END`), one doing I/O and the other doing CPU work?
5. Now, run the same processes, but with the switching behavior set to switch to another process whenever one is **WAITING** for I/O (`-l 1:0,4:100 -c -S SWITCH_ON_IO`). What happens now? Use `-c` and `-p` to confirm that you are right.
6. One other important behavior is what to do when an I/O completes. With `-I IO_RUN_LATER`, when an I/O completes, the process that issued it is not necessarily run right away; rather, whatever was running at the time keeps running. What happens when you run this combination of processes? (`-l 3:0,5:100,5:100,5:100 -S SWITCH_ON_IO -I IO_RUN_LATER -c -p`) Are system resources being effectively utilized?
7. Now run the same processes, but with `-I IO_RUN_IMMEDIATE` set, which immediately runs the process that issued the I/O. How does this behavior differ? Why might running a process that just completed an I/O again be a good idea?
8. Now run with some randomly generated processes: `-s 1 -l 3:50,3:50` or `-s 2 -l 3:50,3:50` or `-s 3 -l 3:50,3:50`. See if you can predict how the trace will turn out. What happens when you use the flag `-I IO_RUN_IMMEDIATE` vs. `-I IO_RUN_LATER`? What happens when you use `-S SWITCH_ON_IO` vs. `-S SWITCH_ON_END`?