



香港城市大學
City University of Hong Kong

Tutorial 4: CPU Scheduling

CS3103

Operating Systems

Scheduling: Introduction

- ▶ Workload assumptions:

1. Each job runs for the **same amount of time**.
2. All jobs **arrive** at the same time (Roughly).
3. Once started, each job runs to **completion**.
4. All jobs only use the **CPU** (i.e., they perform no I/O).
5. The **run-time** of each job is known.

Scheduling Metrics

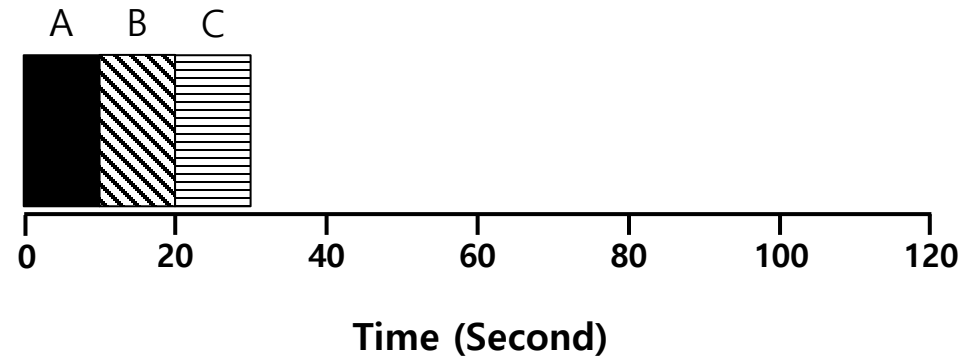
- ▶ Performance metric: **turnaround time**
 - The time at which **the job completes** minus the time at which **the job arrived** in the system.

$$T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$$

- ▶ Another metric is **fairness (response time)**.
 - Performance and fairness are often at odds in scheduling.

First In, First Out (FIFO)

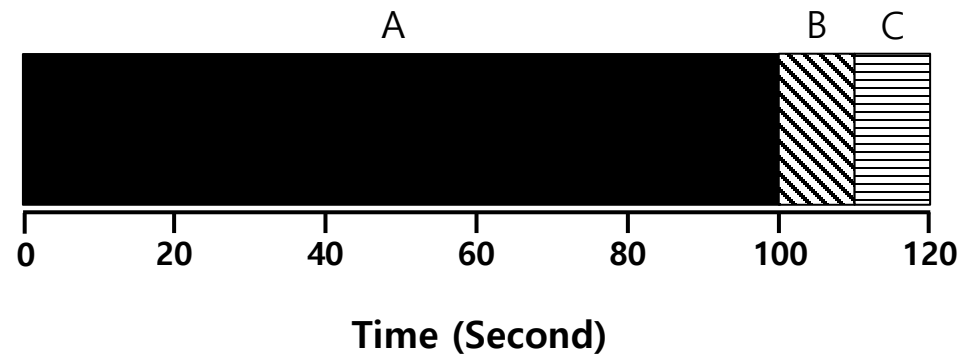
- ▶ First Come, First Served (FCFS)
 - Very simple and easy to implement
- ▶ Example:
 - A arrived just before B which arrived just before C.
 - Each job runs for 10 seconds.



$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

Why FIFO is not that great? – Convoy effect

- ▶ Let's relax assumption 1: Each job **no longer** runs for the same amount of time.
- ▶ Example:
 - A arrived just before B which arrived just before C.
 - A runs for 100 seconds, B and C run for 10 each.

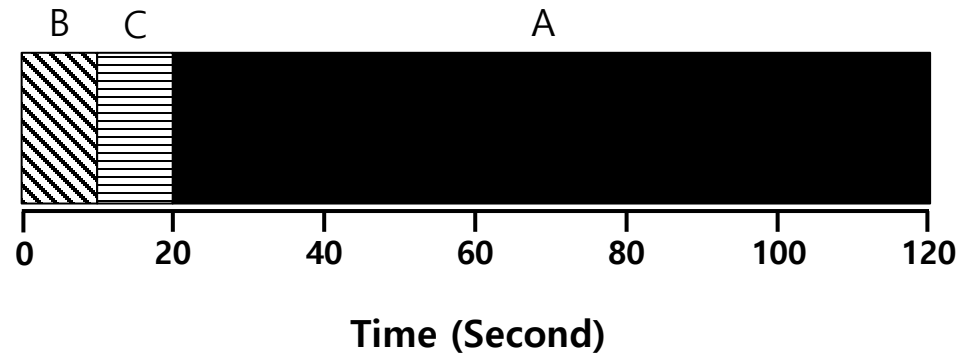


$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = \mathbf{110 \text{ sec}}$$

bloated!

Shortest Job First (SJF)

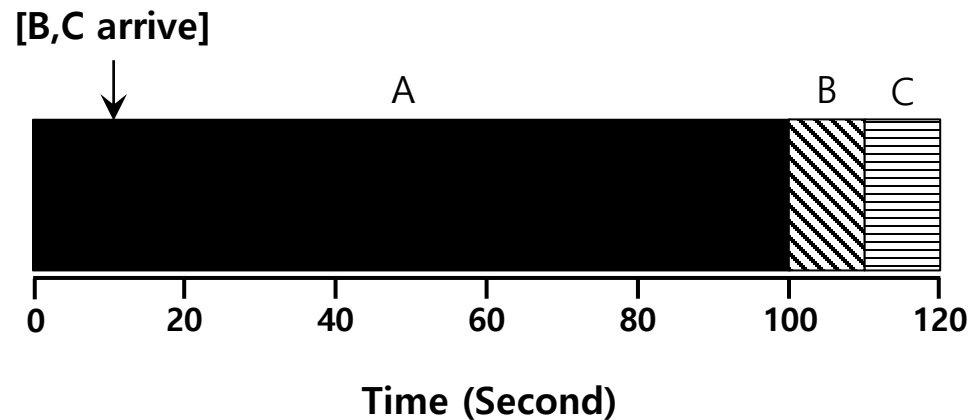
- ▶ Run the shortest job first, then the next shortest, and so on
 - Non-preemptive scheduler
- ▶ Example:
 - A arrived just before B which arrived just before C.
 - A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{10 + 20 + 120}{3} = 50 \text{ sec}$$

SJF with Late Arrivals from B and C (without pre-emptive schedule)

- ▶ Let's relax assumption 2: Jobs can arrive at any time.
- ▶ Example:
 - A arrives at $t=0$ and needs to run for 100 seconds.
 - B and C arrive at $t=10$ and each need to run for 10 seconds



$$\text{Average turnaround time} = \frac{100 + (110 - 10) + (120 - 10)}{3} = 103.33 \text{ sec}$$

Shortest Time-to-Completion First (STCF)

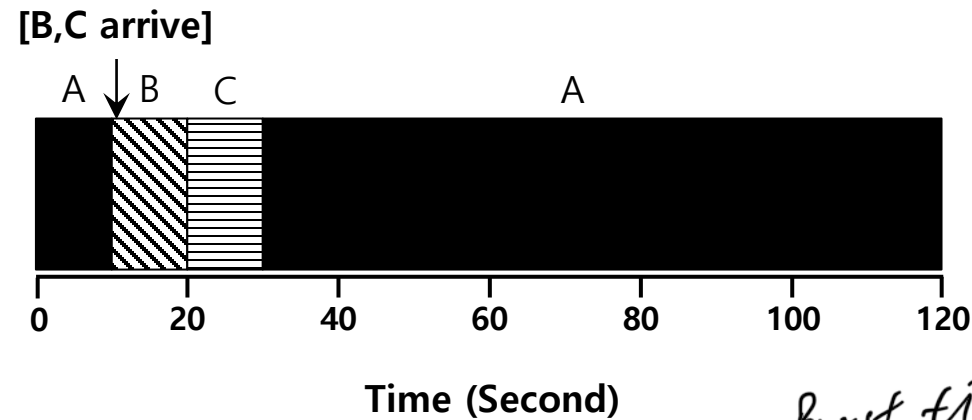
- ▶ Let's relax assumption 3: A job can preempt other jobs.
- ▶ Add **preemption** to SJF
 - Also known as Preemptive Shortest Job First (PSJF)
- ▶ A new job enters the system:
 - Determine of the remaining jobs and new job
 - Schedule the job which has the least time left

Shortest Time-to-Completion First (STCF)

► Example:

- A arrives at $t=0$ and needs to run for 100 seconds.
- B and C arrive at $t=10$ and each need to run for 10 seconds

A → 90s → B, C First → Remaining A



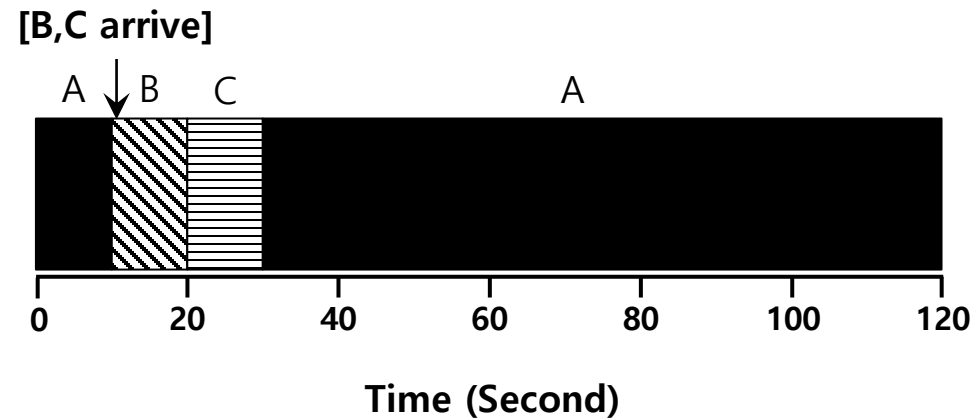
Burst time

$$\text{Average turnaround time} = \frac{(120 - 0) + (20 - 10) + (30 - 10)}{3} = 50 \text{ sec}$$

New scheduling metric: Response time

- ▶ The time from **when the job arrives** to the **first time it is scheduled**.

$$T_{response} = T_{firstrun} - T_{arrival}$$



- STCF and related scheduling policies are not particularly good for response time.

How can we build a scheduler that is
sensitive to response time?

Round Robin (RR) Scheduling

▶ Time slicing Scheduling

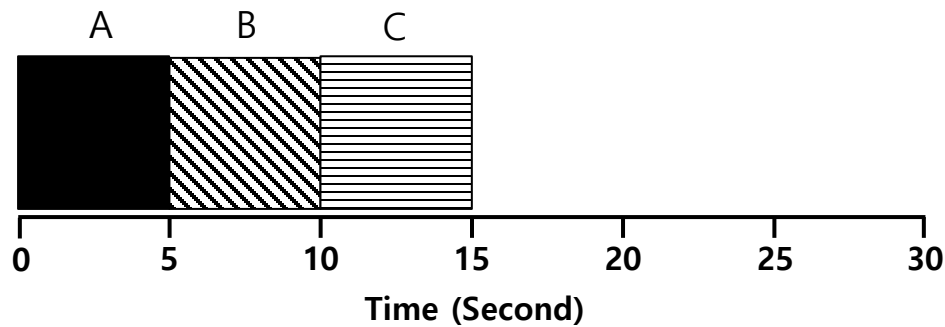
- Run a job for a **time slice** and then switch to the next job in the **run queue** until the jobs are finished.
 - Time slice is sometimes called a scheduling quantum.
- It repeatedly does so until the jobs are finished.
- The length of a time slice must be a *multiple* of the timer-interrupt period.

fairness ✓ performance ↓

**RR is fair, but performs poorly on metrics
such as turnaround time**

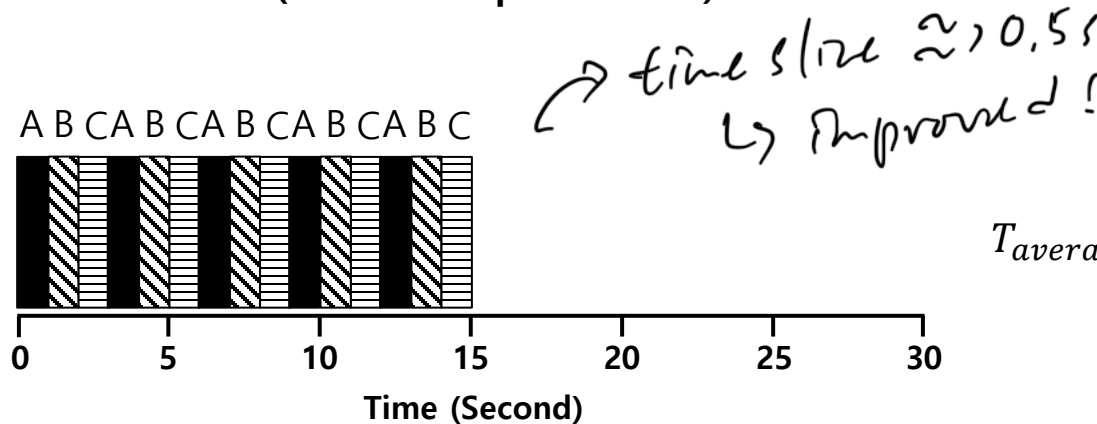
RR Scheduling Example

- ▶ A, B and C arrive at the same time.
- ▶ They each wish to run for 5 seconds.



$$T_{average\ response} = \frac{0 + 5 + 10}{3} = 5sec$$

SJF (Bad for Response Time)



$$T_{average\ response} = \frac{0 + 1 + 2}{3} = 1sec$$

RR with a time-slice of 1sec (Good for Response Time)

The length of the time slice is critical.

- ▶ The shorter time slice
 - Better response time
 - The cost of context switching will dominate overall performance.
- ▶ The longer time slice
 - Amortize the cost of switching
 - Worse response time

**Deciding on the length of the time slice presents
a **trade-off** to a system designer**

Incorporating I/O

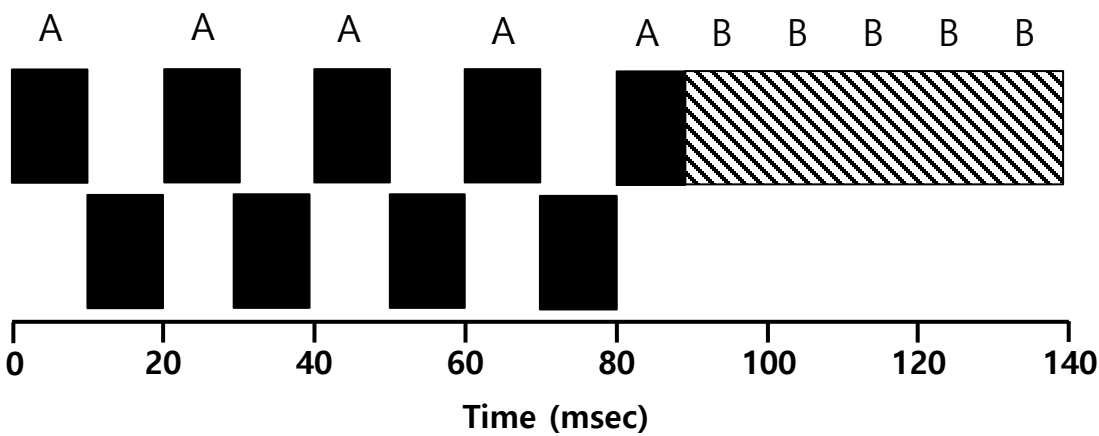
- ▶ Let's relax assumption 4: All programs perform I/O
- ▶ When a job initiates an I/O request.
 - The job is blocked waiting for I/O completion.
 - The scheduler should schedule another job on the CPU.
- ▶ When the I/O completes
 - An interrupt is raised.
 - The OS moves the process from blocked back to the ready state.

Incorporating I/O (Cont.)

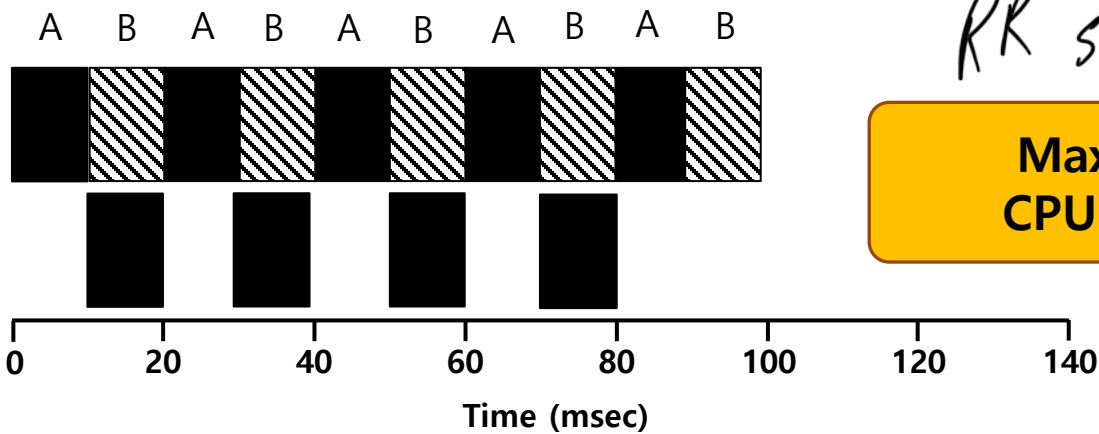
► Example:

- A and B need 50ms of CPU time each.
- A runs for 10ms and then issues an I/O request
 - I/Os each take 10ms
- B simply uses the CPU for 50ms and performs no I/O
- The scheduler runs A first, then B after

Incorporating I/O (Cont.)



Poor Use of Resources



Overlap Allows Better Use of Resources

RR scheduling

Maximize the CPU utilization

Scheduling without *a priori* knowledge

- ▶ Let's relax assumption 5: the OS usually knows very little about the length of each job.
 - Use the recent past to predict the future
 - multi-level feedback queue (covered in the lecture).