



香港城市大學
City University of Hong Kong

Tutorial 5: Deadlocks

CS3103


Operating Systems

Deadlock

Thread 1:	Thread 2:
lock(L1);	lock(L2);
lock(L2);	lock(L1);

► The presence of a cycle

- Thread1 is holding a lock L1 and waiting for another one, L2
- Thread2 that holds lock L2 is waiting for L1 to be release



```
void vector_add(vector_t *v_dst, vector_t *v_src) {  
    Pthread_mutex_lock(&v_dst->lock);  
    Pthread_mutex_lock(&v_src->lock);  
    int i;  
    for (i = 0; i < VECTOR_SIZE; i++) {  
        v_dst->values[i] = v_dst->values[i] + v_src->values[i];  
    }  
    Pthread_mutex_unlock(&v_dst->lock);  
    Pthread_mutex_unlock(&v_src->lock);  
}
```

Why Do Deadlocks Occur?

- ▶ Reason 1:
 - In large code bases, **complex dependencies** arise between components
- ▶ Reason 2:
 - Due to the nature of **encapsulation**
 - Hide details of implementations and make software easier to build in a modular way
 - Such **modularity** *does not mesh* well with locking

Why Do Deadlocks Occur?

- ▶ **Example:** Java Vector class and the method `AddAll()`

```
1  Vector v1, v2;  
2  v1.AddAll(v2);
```

- **Locks** for both the vector being added to (`v1`) and the parameter (`v2`) *need to be acquired*.
 - The routine acquires said locks in some arbitrary order (`v1` then `v2`).
 - If some other thread calls `v2.AddAll(v1)` at nearly the same time → We have the potential for **deadlock**.

Conditional for Deadlock

- ▶ Four conditions need to hold for a deadlock to occur.

Condition	Description
Mutual Exclusion	Threads claim exclusive control of resources that they require.
Hold-and-wait	Threads hold resources allocated to them while waiting for additional resources
No preemption	Resources cannot be forcibly removed from threads that are holding them
Circular wait	There exists a circular chain of threads such that each thread holds one more resources that are being requested by the next thread in the chain

- If any of these four conditions are not met, **deadlock cannot occur**.

Prevention – Circular Wait

- ▶ Provide **a total ordering** on lock acquisition
 - This approach requires *careful design* of global locking strategies.
- ▶ **Example:**
 - There are two locks in the system (L1 and L2)
 - We can prevent deadlock by always acquiring L1 before L2.

Prevention – Hold-and-wait

- Acquire all locks **at once, atomically**.

```
void vector_add(vector_t *v_dst, vector_t *v_src) {  
    // put GLOBAL lock around all lock acquisition...  
    Pthread_mutex_lock(&global);  
    Pthread_mutex_lock(&v_dst->lock);  
    Pthread_mutex_lock(&v_src->lock);  
    Pthread_mutex_unlock(&global);  
    int i;  
    for (i = 0; i < VECTOR_SIZE; i++) {  
        v_dst->values[i] = v_dst->values[i] + v_src->values[i];  
    }  
    Pthread_mutex_unlock(&v_dst->lock);  
    Pthread_mutex_unlock(&v_src->lock);  
}
```

vector-avoid-hold-and-wait.c

- This code guarantees that **no untimely thread switch can occur in the midst of** lock acquisition.
- **Problem:**
 - Require us to know when calling a routine exactly which locks must be held and to acquire them ahead of time.
 - Decrease *concurrency*

Prevention – No Preemption

- ▶ **Multiple lock acquisition** often gets us into trouble because when waiting for one lock **we are holding another**.
- ▶ `trylock()`
 - Used to build a *deadlock-free, ordering-robust* lock acquisition protocol.
 - Grab the lock (if it is available).
 - Or, return -1: you should try again later.

```
void vector_add(vector_t *v_dst, vector_t *v_src) {  
    ↑ top:  
    if (pthread_mutex_trylock(&v_dst->lock) != 0) {  
        goto ↑ top;  
    }  
    if (pthread_mutex_trylock(&v_src->lock) != 0) {  
        retry++;  
        Pthread_mutex_unlock(&v_dst->lock);  
        goto ↑ top;  
    }  
    int i;  
    for (i = 0; i < VECTOR_SIZE; i++) {  
        v_dst->values[i] = v_dst->values[i] + v_src->values[i];  
    }  
    Pthread_mutex_unlock(&v_dst->lock);  
    Pthread_mutex_unlock(&v_src->lock);  
}
```

vector-try-wait.c



Prevention – No Preemption (Cont.)

▶ livelock

- Both systems are running through the code sequence *over and over again*.
- Progress is not being made.
- Solution:
 - Add **a random delay** before looping back and trying the entire thing over again.

Prevention – Mutual Exclusion

► wait-free

- Using powerful **hardware instruction**.
- You can build data structures in a manner that *does not require explicit locking*.

```
1  int CompareAndSwap(int *address, int expected, int new) {  
2      if(*address == expected) {  
3          *address = new;  
4          return 1; // success  
5      }  
6      return 0;  
7  }
```

Prevention – Mutual Exclusion (Cont.)

- ▶ We now wanted to **atomically increment** a value by a certain amount:

```
1  void AtomicIncrement(int *value, int amount) {  
2      do {  
3          int old = *value;  
4      } while ( CompareAndSwap(value, old, old+amount) == 0 );  
5  }
```

- Repeatedly tries to update the value to *the new amount* and uses the `compare-and-swap` to do so.
- **No lock** is acquired
- **No deadlock** can arise
- **livelock** is still a possibility.

Prevention – Mutual Exclusion (Cont.)

► Solution:

- Surrounding this code with a **lock acquire** and **release**.

```
1  void insert(int value) {
2      node_t * n = malloc(sizeof(node_t));
3      assert( n != NULL );
4      n->value = value ;
5      lock(listlock);      // begin critical section
6      n->next    = head;
7      head  = n;
8      unlock(listlock) ;  //end critical section
9  }
```

- **wait-free manner** using the compare-and-swap instruction

```
1  void insert(int value) {
2      node_t *n = malloc(sizeof(node_t));
3      assert(n != NULL);
4      n->value = value;
5      do {
6          n->next = head;
7      } while (CompareAndSwap(&head, n->next, n) == 0);
8  }
```



Deadlock Avoidance via Scheduling

- ▶ In some scenarios **deadlock avoidance** is preferable.
 - **Global knowledge** is required:
 - Which locks various threads might grab during their execution.
 - Subsequently schedules said threads in a way as to guarantee no deadlock can occur.


Example of Deadlock Avoidance via Scheduling (1)

- ▶ We have two processors and four threads.

- Lock acquisition demands of the threads:

	T1	T2	T3	T4
L1	yes	yes	no	no
L2	yes	yes	yes	no

- A smart scheduler could compute that as long as T1 and T2 are not run at the same time, no deadlock could ever arise.




Example of Deadlock Avoidance via Scheduling (2)

- ▶ More contention for the same resources

	T1	T2	T3	T4
L1	yes	yes	yes	no
L2	yes	yes	yes	no

- A possible schedule that guarantees that *no deadlock* could ever occur.



- The total time to complete the jobs is lengthened considerably.

Detect and Recover

- ▶ **Allow deadlock** to occasionally occur and then *take some action*.
 - **Example:** if an OS froze, you would reboot it.
- ▶ Many database systems employ *deadlock detection and recovery technique*.
 - A deadlock detector **runs periodically**.
 - Building a **resource graph** and checking it for cycles.
 - In deadlock, the system **need to be restarted**.