



香港城市大學
City University of Hong Kong

Tutorial 7: Page Replacement Policy

CS3103

Operating Systems

Beyond Physical Memory: Policies

- ▶ Memory pressure forces the OS to start **paging out** pages to make room for actively-used pages.
- ▶ Deciding which page to evict is encapsulated within the replacement policy of the OS.

Cache Management

- ▶ Goal in picking a replacement policy for this cache is to minimize the number of cache misses.
- ▶ The number of cache hits and misses let us calculate the *average memory access time(AMAT)*.

$$AMAT = (P_{Hit} * T_M) + (P_{Miss} * T_D)$$

Argument	Meaning
T_M	The cost of accessing memory
T_D	The cost of accessing disk
P_{Hit}	The probability of finding the data item in the cache(a hit)
P_{Miss}	The probability of not finding the data in the cache(a miss)

The Optimal Replacement Policy

- ▶ Leads to the fewest number of misses overall
 - Replaces the page that will be accessed furthest in the future
 - Resulting in the **fewest-possible** cache misses
- ▶ Serve only as a comparison point, to know how close we are to **perfect**

Tracing the Optimal Policy

Reference Row

0 1 2 0 1 3 0 3 1 2 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	2	0,1,3
0	Hit		0,1,3
3	Hit		0,1,3
1	Hit		0,1,3
2	Miss	3	0,1,2
1	Hit		0,1,2

Hit rate is $\frac{Hits}{Hits+Misses} = 54.6\%$

Future is not known.

A Simple Policy: FIFO

- ▶ Pages were placed in a queue when they enter the system.
- ▶ When a replacement occurs, the page on the tail of the queue(the “**First-in**” pages) is evicted.
 - It is simple to implement, but can’t determine the importance of blocks.

Tracing the FIFO Policy

Reference Row

0 1 2 0 1 3 0 3 1 2 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	0	1,2,3
0	Miss	1	2,3,0
3	Hit		2,3,0
1	Miss		3,0,1
2	Miss	3	0,1,2
1	Hit		0,1,2


Hit rate is $\frac{\text{Hits}}{\text{Hits} + \text{Misses}} = 36.4\%$

Even though page 0 had been accessed a number of times, **FIFO still kicks it out.**

BELADY'S ANOMALY

- ▶ We would expect the cache hit rate to **increase** when the cache gets larger. But in this case, with FIFO, it gets worse.

Reference Row											
1	2	3	4	1	2	5	1	2	3	4	5




Another Simple Policy: Random

- ▶ Picks a random page to replace under memory pressure.
 - It doesn't really try to be too intelligent in picking which blocks to evict.
 - Random does depends entirely upon how lucky Random gets in its choice.

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	0	1,2,3
0	Miss	1	2,3,0
3	Hit		2,3,0
1	Miss	3	2,0,1
2	Hit		2,0,1
1	Hit		2,0,1

Random Performance

- Sometimes, **Random is as good as optimal**, achieving 6 hits on the example trace.



Using History

- ▶ Lean on the past and use **history**.
 - Two type of historical information.

Historical Information	Meaning	Algorithms
recency	The more recently a page has been accessed, the more likely it will be accessed again	LRU
frequency	If a page has been accessed many times, It should not be replaced as it clearly has some value	LFU

Using History : LRU

- Replaces the least-recently-used page.


Reference Row

0 1 2 0 1 3 0 3 1 2 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		1,2,0
1	Hit		2,0,1
3	Miss	2	0,1,3
0	Hit		1,3,0
3	Hit		1,0,3
1	Hit		0,3,1
2	Miss	0	3,1,2
1	Hit		3,2,1

Workload Example : The No-Locality Workload


- ▶ Each reference is to a random page within the set of accessed pages.
 - Workload accesses 100 unique pages over time.
 - Choosing the next page to refer to at random



When the cache is large enough to fit the entire workload, it also **doesn't matter** which policy you use.

Workload Example : The 80-20 Workload


- ▶ Exhibits locality: 80% of the **references** are made to 20% of the page
- ▶ The remaining 20% of the **references** are made to the remaining 80% of the pages.



LRU is more likely to hold onto the **hot pages**.

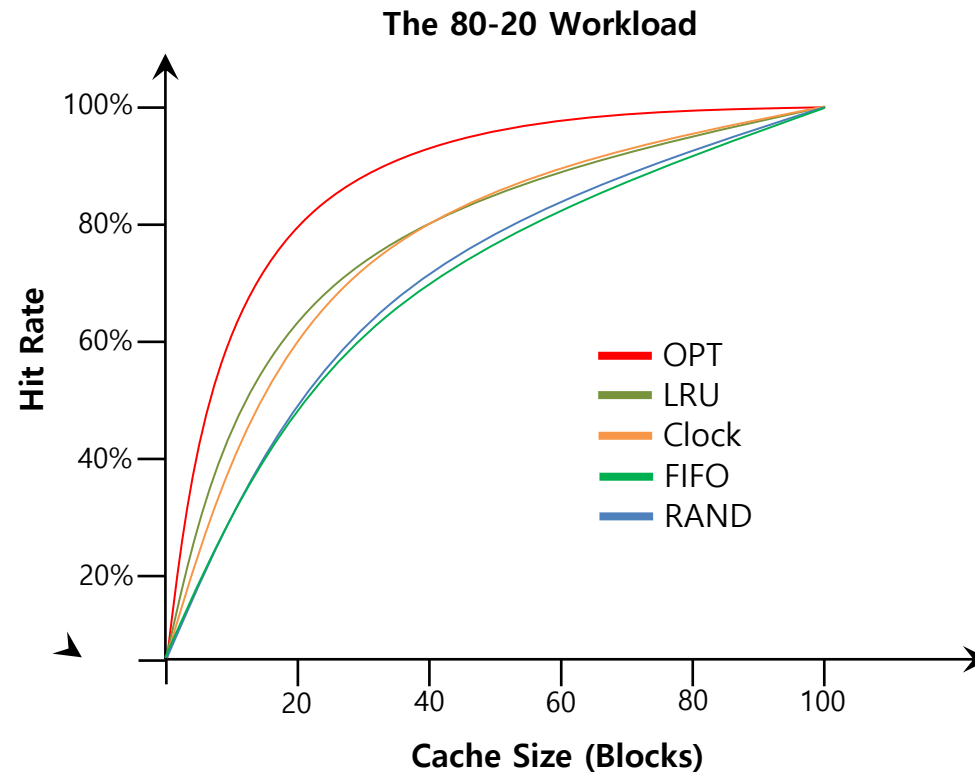
Workload Example : The Looping Sequential

- ▶ Refer to 50 pages in sequence.
 - Starting at 0, then 1, ... up to page 49, and then we Loop, repeating those accesses, for total of 10,000 accesses to 50 unique pages.



Workload with Clock Algorithm

- ▶ Clock algorithm doesn't do as well as perfect LRU, it does better than approach that don't consider history at all.



Considering Dirty Pages



- ▶ The hardware include a **modified bit** (a.k.a **dirty bit**)
 - Page has been **modified** and is thus **dirty**, it must be written back to disk to evict it.
 - Page has not been modified; the eviction is free.

Page Selection Policy

- ▶ The OS has to decide when to bring a page into memory.
- ▶ Presents the OS with some different options.

Prefetching

- ▶ The OS guess that a page is about to be used, and thus bring it in ahead of time.



Secondary
Storage

Page 2 likely soon be accessed and thus should be brought into memory too

Thrashing

- ▶ Memory is **oversubscribed** and the memory demands of the set of running processes **exceeds** the available physical memory.
 - Decide not to run a subset of processes.
 - Reduced set of processes working sets fit in memory.

